

Kernel and processes

Signals

Signals/ Events

- Process interact with kernel using system calls
- In computer system special events require the kernel to interact with processes
- In UNIX/LINUX these events are called signals
 - Exceptions
 - illegal memory access, division by zero, ...
 - User generated
 - process abortion , Ctrl-C
 - generated by system calls
 - son process termination, timer, kill

How to handle signals

- Polling
 - All processes poll the kernel to existing events
 - Not efficient
 - Not correct
 - programmer should implement this “system”feature”
- kernel can manage notification (and scheduling)
 - of processes requesting notification

Upcall: User-level event delivery

- Notify user process of some event that needs to be handled right away
 - Time expiration
 - Real-time user interface
 - Time-slice for user-level thread manager
 - Interrupt delivery for VM player
 - Asynchronous I/O completion (async/await)
- UNIX signal


Upcalls vs Interrupts

- Signal handlers = interrupt vector
- Signal stack = interrupt stack
- Automatic save/restore registers = transparent resume
- Signal masking: signals disabled while in signal handler


Upcall: Before

- Process previously register a signal handler
 - will execute in user level
 - kernel will redirect execution to it
- Kernel was notified of the event
 - verifies that process registered and handler

```
...  
x = y + z; ←  
...  
Program Counter
```



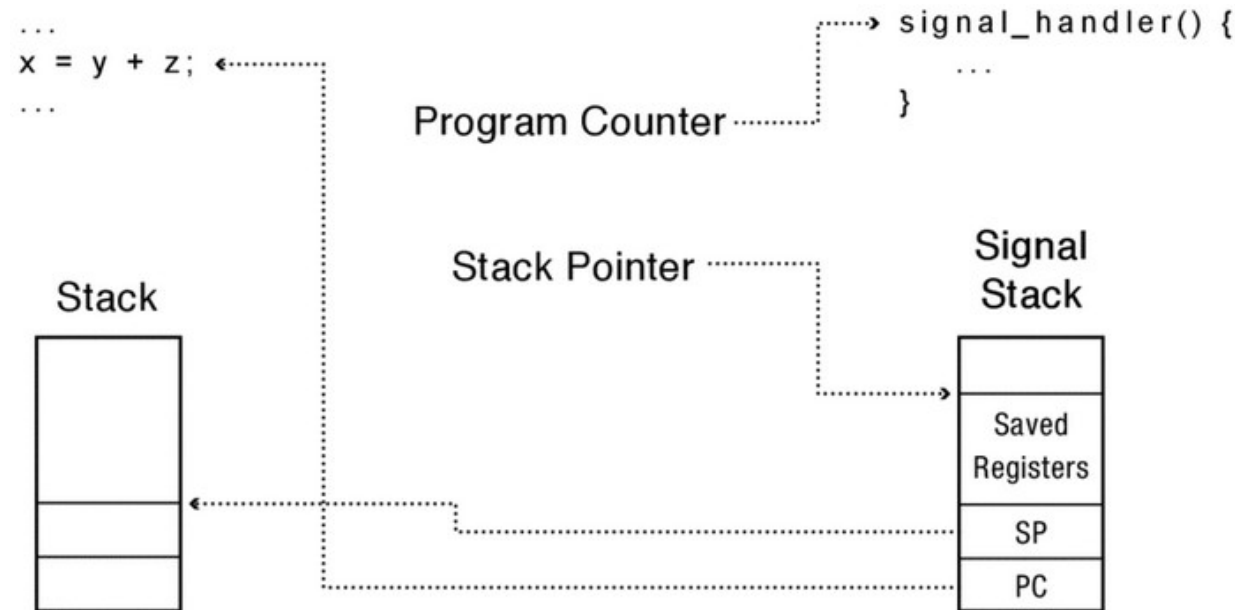
```
signal_handler() {  
    ...  
}
```



Upcall: During

- In kernel mode
 - created signal stack
 - correct signal return
 - updated SP and PC
 - for process to execute handler

- In user mode
 - execution of handler
 - Return to resume



Signal handling in UNIX

- Default handling
 - Kernel handles the signal
 - Process is not notified
 - Most of times it is terminated
- Ignore the signal
- handle the event with a user level handler
- Two signals can not be handled or ignored:
 - SIGKILL
 - SIGSTOP

Unix signals

Signal	Code	Default action	Cause
SIGHUP	1	Terminates process	Terminal line hangup
SIGINT	2	Terminates process	interrupt program (CTRL-C)
SIGQUIT	3	Terminates process + core dumps	Quit program
SIGKILL	9	Terminates process	Kill program
SIGALRM	14	Terminates process	Timer expired
SIGCHLD	20	Ignored	child status has changed
SIGSTOP	17	Stop process	Stop signal generated from keyboard (CTRL-Z)
SIGSCONT	19	Continue if stoped	Continue after stop
SIGUSR1	31	Terminates process + core dumps	User defined signal 1

Unix signals

Table 4-1. UNIX signals

Signal	Description	Default Action	Available In	Notes
SIGABRT	process aborted	abort	APSB	
SIGALRM	real-time alarm	exit	OPSB	
SIGBUS	bus error	abort	OSB	6
SIGCHLD	child died or suspended	ignore	OJSB	4
SIGCONT	resume suspended process	continue/ignore	JSB	
SIGEMT	emulator trap	abort	OSB	
SIGFPE	arithmetic fault	abort	OAPSB	
SIGHUP	hang-up	exit	OPSB	2
SIGILL	illegal instruction	abort	OAPSB	
SIGINFO	status request (control-T)	ignore	B	
SIGINT	tty interrupt (control-C)	exit	OAPSB	3
SIGIO	async I/O event	exit/ignore	SB	
SIGIOT	I/O trap	abort	OSB	1
SIGKILL	kill process	exit	OPSB	
SIGPIPE	write to pipe with no readers	exit	OPSB	
SIGPOLL	pollable event	exit	S	
SIGPROF	profiling timer	exit	SB	
SIGPMR	power fail	ignore	OS	
SIGQUIT	tty quit signal (control-\)	abort	OPSB	
SIGSEGV	segmentation fault	abort	OAPSB	1
SIGSTOP	stop process	stop	JSB	
SIGSYS	invalid system call	exit	OAPSB	
SIGTERM	terminate process	exit	OAPSB	2
SIGTRAP	hardware fault	abort	OSB	
SIGTSTP	tty stop signal (control-Z)	stop	JSB	
SIGTTIN	tty read from background process	stop	JSB	5
SIGTTOU	tty write from background process	stop	JSB	
SIGURG	urgent event on I/O channel	ignore	SB	
SIGUSR1	user-definable	exit	OPSB	
SIGUSR2	user-definable	exit	OPSB	
SIGVTALRM	virtual time alarm	exit	SB	
SIGWINCH	window size change	ignore	SB	
SIGXCPU	exceed CPU limit	abort	SB	
SIGXFSZ	exceed file size limit	abort	SB	

Availability: O Original SVR2 signal
 B 4.3 BSD
 P 1 POSIX.1
 1 cannot be caught, blocked, or ignored.
 2 Not reset to default, even in System V implementations.
 3 Default action is to exit in SVR4, ignore in 4.3BSD.
 4 Process can choose to allow background writes without generating this signal.
 5 Called SIGCLD in SVR3 and earlier releases.
 6 Called SIGCLD in SVR3 and earlier releases.

Notes: A ANSI C
 S SVR4
 J POSIX.1, only if job control is supported

Signal Generation

- Exceptions
 - When an exception occurs in the process (for instance, an attempt to execute an illegal instruction), the kernel notifies the process by sending it a signal.
- Other processes
 - A process may send a signal to another process, or set of processes, through the **kill** or **sigsend** system calls.
 - A process may even send a signal to itself.
- Terminal Interrupts
 - Certain keyboard characters, such as control-C or control-\, send signals to the foreground process on that terminal.
 - The **stty** command allows the user to bind each terminal-generated signal to a specific key.
- Job control
 - Background processes that try to read or write to the terminal are sent job control signals.
 - Job control shells such as **csh** and **ksh** use signals to manipulate foreground and background processes.
 - When a process terminates or is suspended, the kernel notifies its parent via a signal.

Signal Generation

- Quotas
 - When a process exceeds its CPU or file size limits, the kernel sends a signal to the process.
- Notifications
 - A process may request notification of certain events, such as a device being ready for I/O.
 - The kernel informs the process via a signal.
- Alarms
 - A process may set an alarm for a certain time;
 - when it expires, the kernel notifies the process through a signal

Signal handler

- `void (*sighandler_t)(int);`
- Registration of signal handler
 - `#include <signal.h>`
 - `sighandler_t signal(int, sighandler_t);`
 - 1st parameter
 - signal code (SIGXXX)
 - 2nd parameter
 - pointer to function
 - or SIG_IGN
 - or SIG_DFL
- Returns the previous signal handler
- Different OS implement signals in different ways
 - signal handler should be re-assigned

Unreliable Signals

- Bigeste concern
 - Signal delivery

```
int    sig_int(){  
    /* my signal handling function */  
    ...  
    /* re install the handler */  
    signal(SIGINT, sig_int);  
    ...  
}  
  
int main(){  
    signal(SIGINT, sig_int); /* install the handler */  
    ... /* process the signal ... */  
}
```

- To catch new signal occurrences
 - Users must re-install it
- Synchronization problems
 - Signal generated in fast succession?
- No mechanisms to temporary block a signal....

Reliable signals

- Solves previous problems
 - BSD + POSIX
- Persistent handlers
 - Signal handlers remain installed
 - No windows between signal catch and re-install
- Masking
 - A signal can be masked/blocked temporarily
 - Kernel keeps signal until it is unblocked
- Unblock and wait
 - A process can be blocked until it receives a signal

Signal handler

- `int sigaction(int, const struct sigaction * newhandler, struct sigaction * old);`
 - System call to configure system handling
 - Same behavior in all systems
 - 1st parameter : specifies a signal number.
 - 2nd parameter: new handler
 - 3rd parameter: old handler
- Signal handler configuration - struct sigaction
 - `void (*sa_handler)()`
 - signal handler function
 - `void (*sa_sigaction)(int signum, siginfo_t *siginfo, void *uctx);`
 - alternative signal handler function
 - `sigset_t sa_mask`
 - signals that should be blocked during signal handler
 - `int sa_flags`
 - signal handling configuration
 - `void (*sa_restorer)();`

-

User Generated signals

- `int kill(pid_t pid, int sig);`
 - `#include <sys/types.h>`
 - `#include <signal.h>`
 - 1st parameter
 - pid of process to receive signal
 - -1 all processes
 - 2nd parameter
 - Signal number
 - Returns 0 in case of success
 - -1

Signal Management API

- kill
 - sends signal to process
- `typedef void (*sighandler_t)(int);`
 - Signal handler
- `Signal() / sigaction()`
 - sets the disposition of the a signal a handler
 - Can also disable or reset
- If signal is not handled will probably terminate process