

Threads

- Operating Systems Principles & Practice
 - Volume II, chapter 5
- Unix Internals
 - Chapter, 3
- Multithreaded Programming Guide, Sun
 - Chapters 1 and 2

Concurrency

- In the real world different activities often proceed at the same time.
 - But inter-related (on activity affect the others)
 - They are concurrent
- Computers are concurrent
 - Multiple resources
 - dozen processors, 10 disks, and 4 network interfaces; a workstation might have a dozen active I/O devices including a screen, keyboard, mouse, camera, microphone, speaker, wireless network interface, wired network interface, printer, scanner, and disk drive
 - Multiple CPUs / Cores
 - Multiple users/applications



Concurrency

- Relevant to the programmer
 - Network services need to be able to handle multiple requests from their client
 - Most applications today have user interfaces while simultaneously executing application logic
 - Parallel programs need to be able to map work onto multiple processors
 - Need to mask the latency of disk and network operations

Threads

- A Thread is an independent stream of instructions that can be schedule to run as such by the OS.
- Think of a thread as a “procedure” that runs independently from its main program.
- Multi-threaded programs are where several procedures are able to be scheduled to run simultaneously and/or independently by the OS.
- A Thread exists within a process and uses the process resources.


Thread




Thread 1:
DrawScene()



Thread 2:
DrawScene()



Thread 3:
DrawWidgets()



Thread 4:
GetData()

- Threads only duplicate the essential resources it needs to be independently schedulable.
- A thread will die if the parent process dies.
- A thread is “lightweight” because most of the overhead has already been accomplished through the creation of the process.

- Private

- Processor register
- Stack

- Shared

- Memory
- Resources

Why Use Threads


- The primary motivation behind Pthreads
 - is improving program performance.
 - Easing development
- Can be created with much less OS overhead.
- Needs fewer system resources to run.

When to use

- Message multiplexing in the same channel
 - Each thread handles a type of messages
 - Releases main thread to handle other messages
- Multiple channels
 - Each thread handles a channel
- Synchronized wait / Event notification
 - One thread is blocked waiting for an event
- Shared memory between parallel executions

Multithreaded systems

- A system that allows processes with multiple threads
 - Are called multithreaded
- OS can be split into:

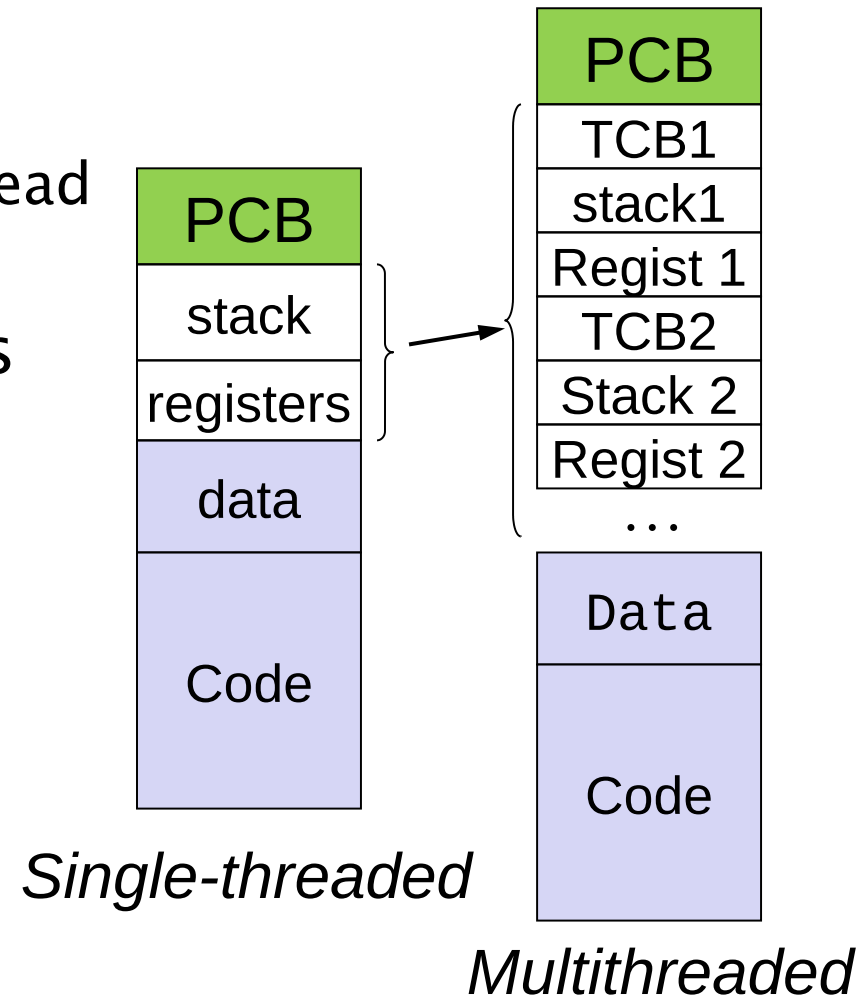


Thread

- Execution thread
 - Execution unit
 - private
 - Processor register
 - Stack
 - Shared (inside the same process)
 - Memory
 - resources
- One process contains one thread when started

Process structure

- With threads the process image changes:
 - Each thread contains a local TCBThread Control Block.
- In a process each thread contains its own stack
 - Local variables are local to each thread
- Threads share
 - Code
 - Global variables
 - Resources (FILES, IPC)



Process vs threads

Processes	Threads
Hierarquical struture <ul style="list-style-type: none">• Parent• Child	Flat stucture <ul style="list-style-type: none">• All siblings
Independent data space <ul style="list-style-type: none">• No shared variables	Shared data space Shared DS Shared heap
Creation os expensive	Creation is cheap <ul style="list-style-type: none">• More than 20 x faster

PLATFORM	fork()			pthread_create()		
	REAL	USER	SYSTEM	REAL	USER	SYSTEM
AMD 2.4 GHz Opteron (8cpus/node)	41.07	60.08	9.01	0.66	0.19	0.43
IBM 1.9 GHz POWER5 p5-575 (8cpus/node)	64.24	30.78	27.68	1.75	0.69	1.1
IBM 1.5 GHz POWER4 (8cpus/node)	104.05	48.64	47.21	2.01	1	1.52
INTEL 2.4 GHz Xeon (2 cpus/node)	54.95	1.54	20.78	1.64	0.67	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.54	1.07	22.22	2.03	1.26	0.67

Multithreaded Programs

- Best used with programs that can be organized
 - into discrete, independent tasks which can execute concurrently.
- Threads can be
 - interchanged,
 - interleaved
 - overlapped

Multithreaded Programs

Programmer's
View

.
.
.
x = x + 1;
y = y + x;
z = x + 5y;
.
.
.

Possible
Execution
#1

.
.
.
x = x + 1;
y = y + x;
z = x + 5y;
.
.
.

Possible
Execution
#2


.
.
.
x = x + 1;
.....
Thread is suspended.
Other thread(s) run.
Thread is resumed.
.....
y = y + x;
z = x + 5y;

Possible
Execution
#3


.
.
.
x = x + 1;
y = y + x;
.....
Thread is suspended.
Other thread(s) run.
Thread is resumed.
.....
z = x + 5y;

Multithreaded Programs


One Execution



Another Execution



Another Execution



Unpredictable
Speed

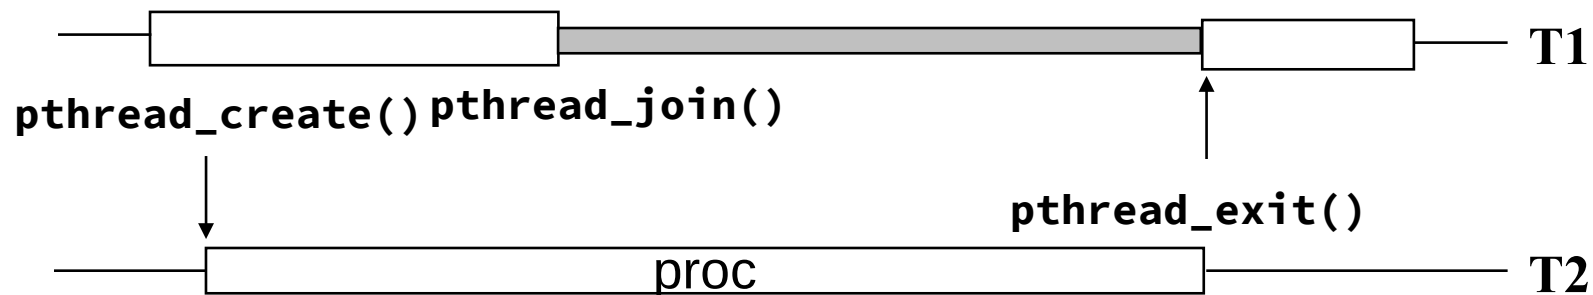
Simple Threads API

- `void thread_create (thread, func, arg)`
 - Create a new thread, storing information about it in `thread`. Concurrently with the calling thread, `thread` executes the function `func` with the argument `arg`.
- `void thread_yield ()`
 - The calling thread voluntarily gives up the processor to let some other thread(s) run. The scheduler can resume running the calling thread whenever it chooses to do so.
- `int thread_join (thread)`
 - Wait for `thread` to finish if it has not already done so; then return the value passed to `thread_exit` by that thread. Note that `thread_join` may be called only once for each thread.
- `void thread_exit (ret)`
 - Finish the current thread. Store the value `ret` in the current thread's data structure. If another thread is already waiting in a call to `thread_join`, resume it.

Main Thread

Thread2

```
pthread_create(&thread2, NULL, proc, &arg);  
pthread_join(thread2, &status);  
pthread_exit(&status); }
```



- There are several APIs
 - Win32 threads.
 - C-Threads (user level)
 - Pthreads
 - POSIX IEEE 1003.1c, published in 1995
- POSIX defines functions for the management of threads
 - Functions/data started with the prefix **pthread_**
- Definitions available in the **pthread.h** file
- Code should be linked with the pthread library
 - **-lpthread**

Thread Identification

- Each thread has a unique identifier of type `pthread_t`
- A thread knows its ID calling **`pthread_self()`**
 - `pthread_t pthread_self();`
- To compare thread identifiers use
 - `int pthread_equal(pthread_t, pthread_t)`
- To print use format `%lu` (long unsigned)

Thread creation

- The `main()` method comprises a single, default thread.
- `pthread_create()` creates a new thread and makes it executable.
- The maximum number of threads that may be created by a process in implementation dependent.
- Once created, threads are peers, and may create other threads.

Thread creation

- A thread is started with

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void *(*start_routine) (void *),  
                  void *arg);
```

- 1st parameter – Pointer to thread identifier (out)
- 2nd parameter – Pointer to thread attributes (IN)
 - Can be NULL .
- 3rd parameter – Pointer to function containing the thread code
 - Function should be: **void * (func*) (void * arg).**
- 4th parameter – Pointer to thread arguments
 - Pointer to array, structure, int, (can be NULL)
- Returns 0 if successful

Data transfer

- Data can be transmitted into the thread in several ways
 - Global variables
 - Accessible by all threads (synchronization should be applied)
 - 4th parameter of `pthread_create`
 - This parameter points to any data structure the programmer defines
 - Not use same memory location to multiple threads
 - Coherency not guaranteed
- Out data follows similar pattern

Thread termination

- Several ways to terminate a thread:
 - The thread is complete and returns
 - The `pthread_exit()` method is called
 - The `pthread_cancel()` method is invoked
 - The `exit()` method is called
- The `pthread_exit()` routine is called
 - By the exiting thread
 - after it has completed its work and it no longer is required to exist.

- The `pthread_cancel()` routine is called
 - By any thread
 - Terminates other running thread
- If the main thread finishes with `pthread_exit`
 - the other threads will continue to exist
- The `pthread_exit()` method does not close files;
 - any files opened inside the thread will remain open, so cleanup must be kept in mind.
- `Exit()` in `main()` will terminate all threads

- A thread kills itself by calling
 - `int pthread_exit(void *ret);`
 - 1st parameter – pointer to the return code/data
 - Can be pointer to any data type
 - Memory location should be accessible outside (either global variable or malloc)
 - RETURN VALUE – This function does not return to the caller.
 - ERRORS – This function always succeeds.
- `fazer return(cod)` implicitly calls `pthread_exit()`
- After `pthread_exit` resources are maintained
 - Resources are released only after `pthread_join()`.
- `int pthread_detach(pthread_t);`
 - Resources are immediately released
 - `pthread_join()` can not be done.

Wait for a thread

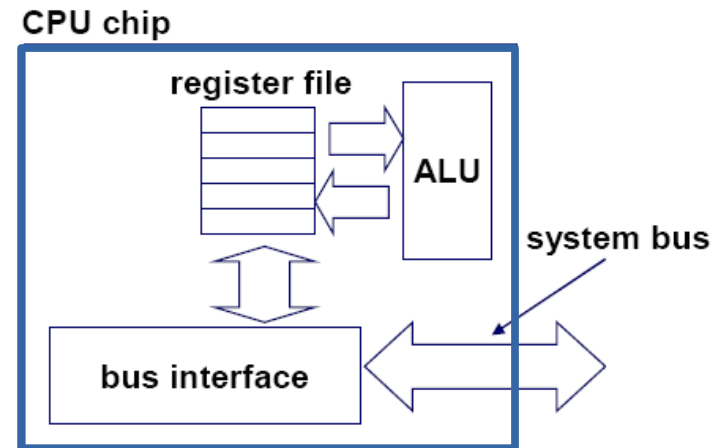
- A thread waits for another thread executing **pthread_join**
 - To release resources
 - To fetch returned data
 - **int pthread_join(pthread_t thread, void **retval);**
 - 1st parameter
 - thread identifier.
 - 2nd parameter
 - Pointer to location of returned value
- function waits for the thread specified by thread to terminate.
- If that thread has already terminated, then pthread_join() returns immediately.
- Only one thread can wait/join another thread

Multicore CPUs

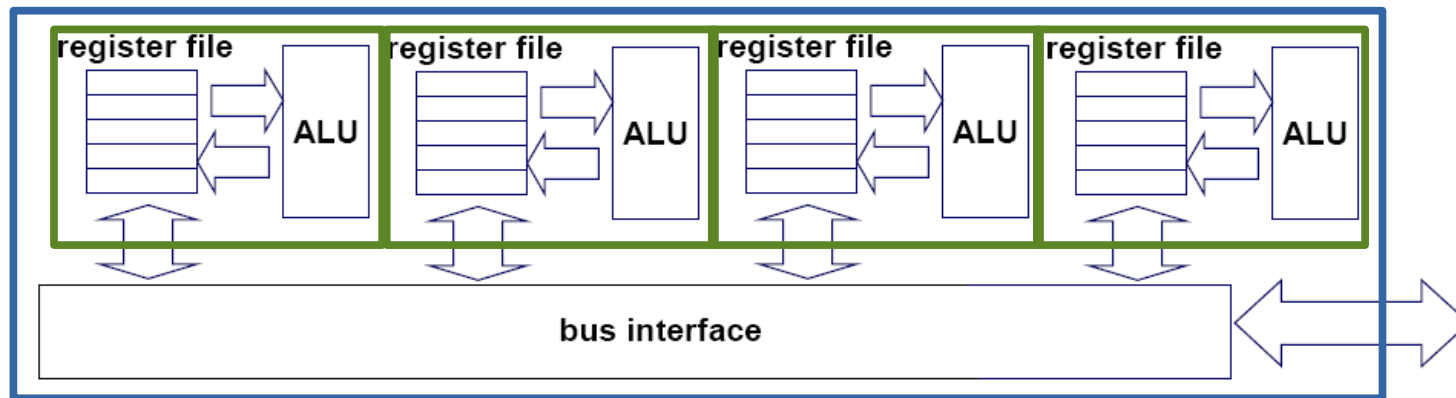
- A core is a processing unit
 - Inside the same packaging (“CPU”)
 - A dual/quad core CPU include 2/4 cores
- Evolution from multi-processor
 - Several CPU connected by a BUS
- First dualcore CPU
 - Power 4, from IBM in 2000.
- First Dual core from intel
 - Jan 2006.
- Commodity processors <12 cores
- Specialized processors < 100 cores

Multicore Architecture

- Single Core Processors



- Multi-core Processors



Multicore Architecture

- Cores share RAM memory
- Caches are:
 - L1 – private (one per cores)
 - L2 – private on most systems
 - L3 – shared
- Threads are assigned to cores
 - Possibly different cores
 - Can be controlled with thread affinity

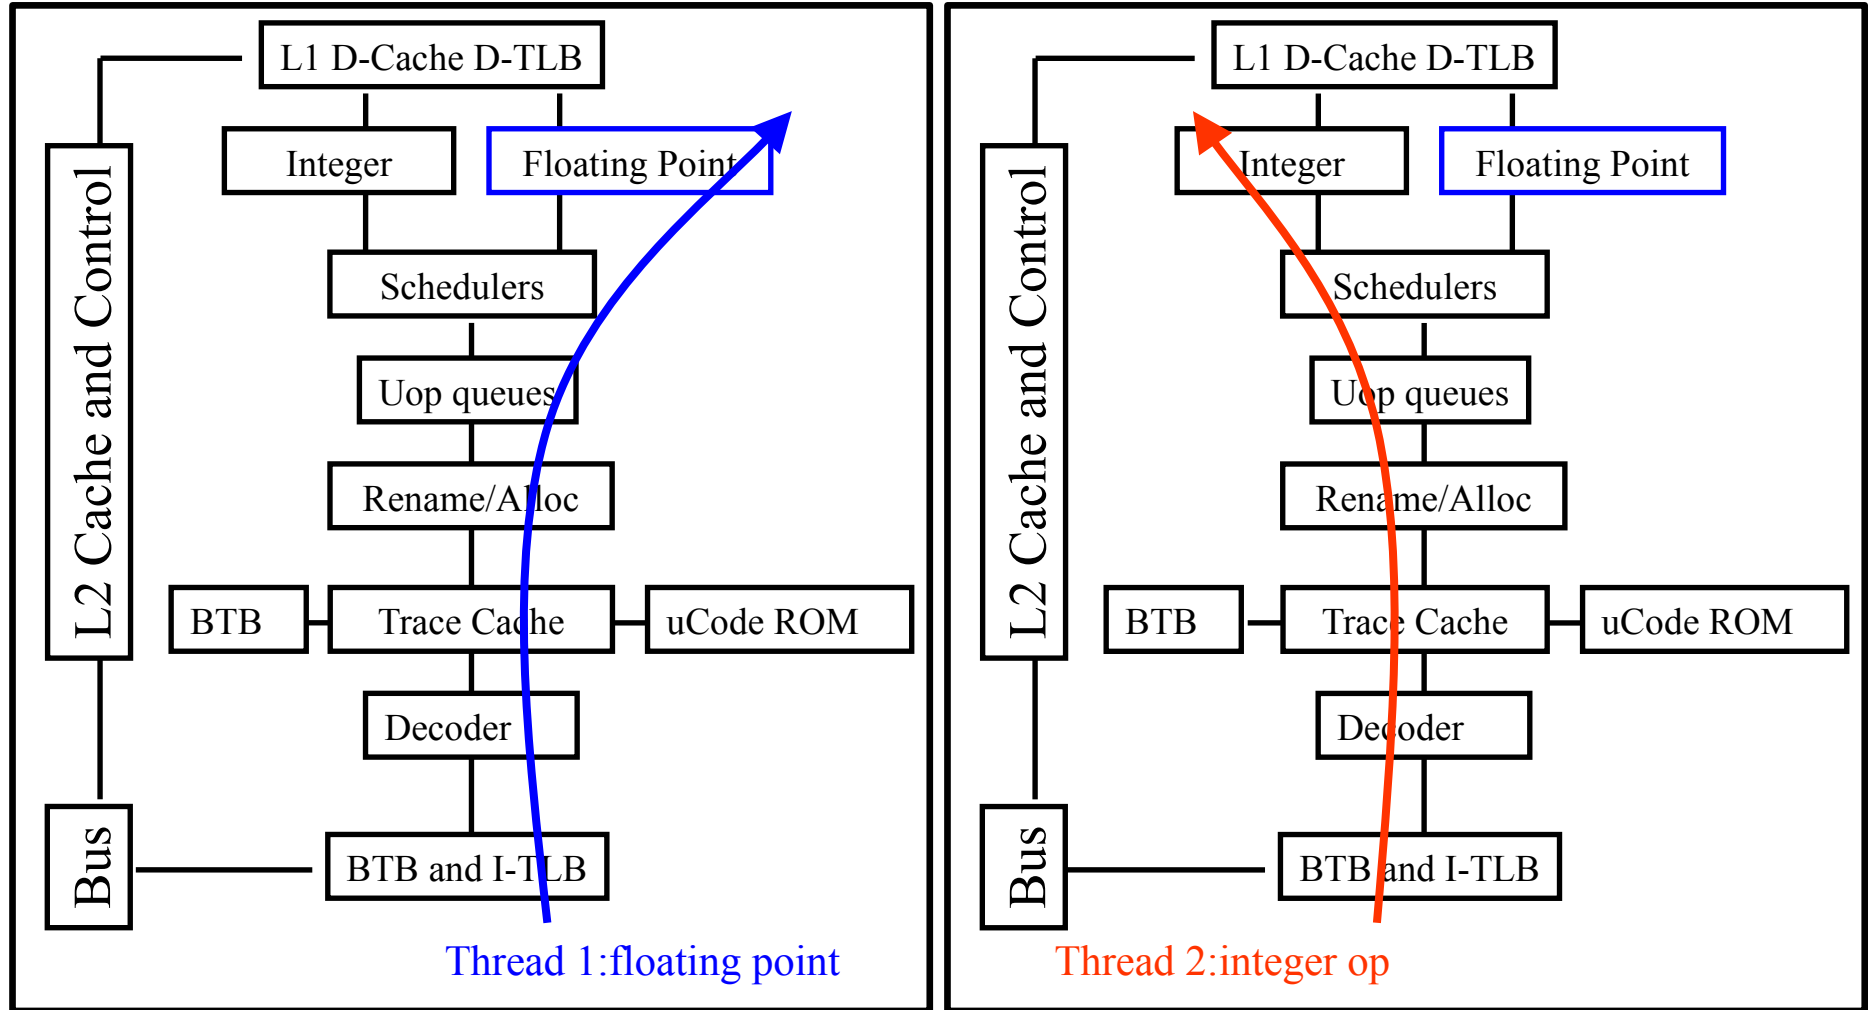
Cache coherence

- Since each core has one private cache
 - Data coherence is fundamental
- `Int var_i = 125` ← shared by thread 1 and 2
 - Thread 1 runs on core 1
 - Thread 2 runs on core 2
- `var_i = 0`
- `var_i` is invalidated
- `Priv = var_i`
 - Cache miss
 - `var_i` copied from
 - From thread 1 cache

Cache coherence

- Since each core has one private cache
 - Data coherence is fundamental
- `Int var_i, var_j = 125` ← shared by thread 1 and 2
 - Thread 1 runs on core 1
 - Thread 2 runs on core 2
- `var_i = 0`
- `var_i` is invalidated
- `Priv = var_j`
 - Cache miss or hit?

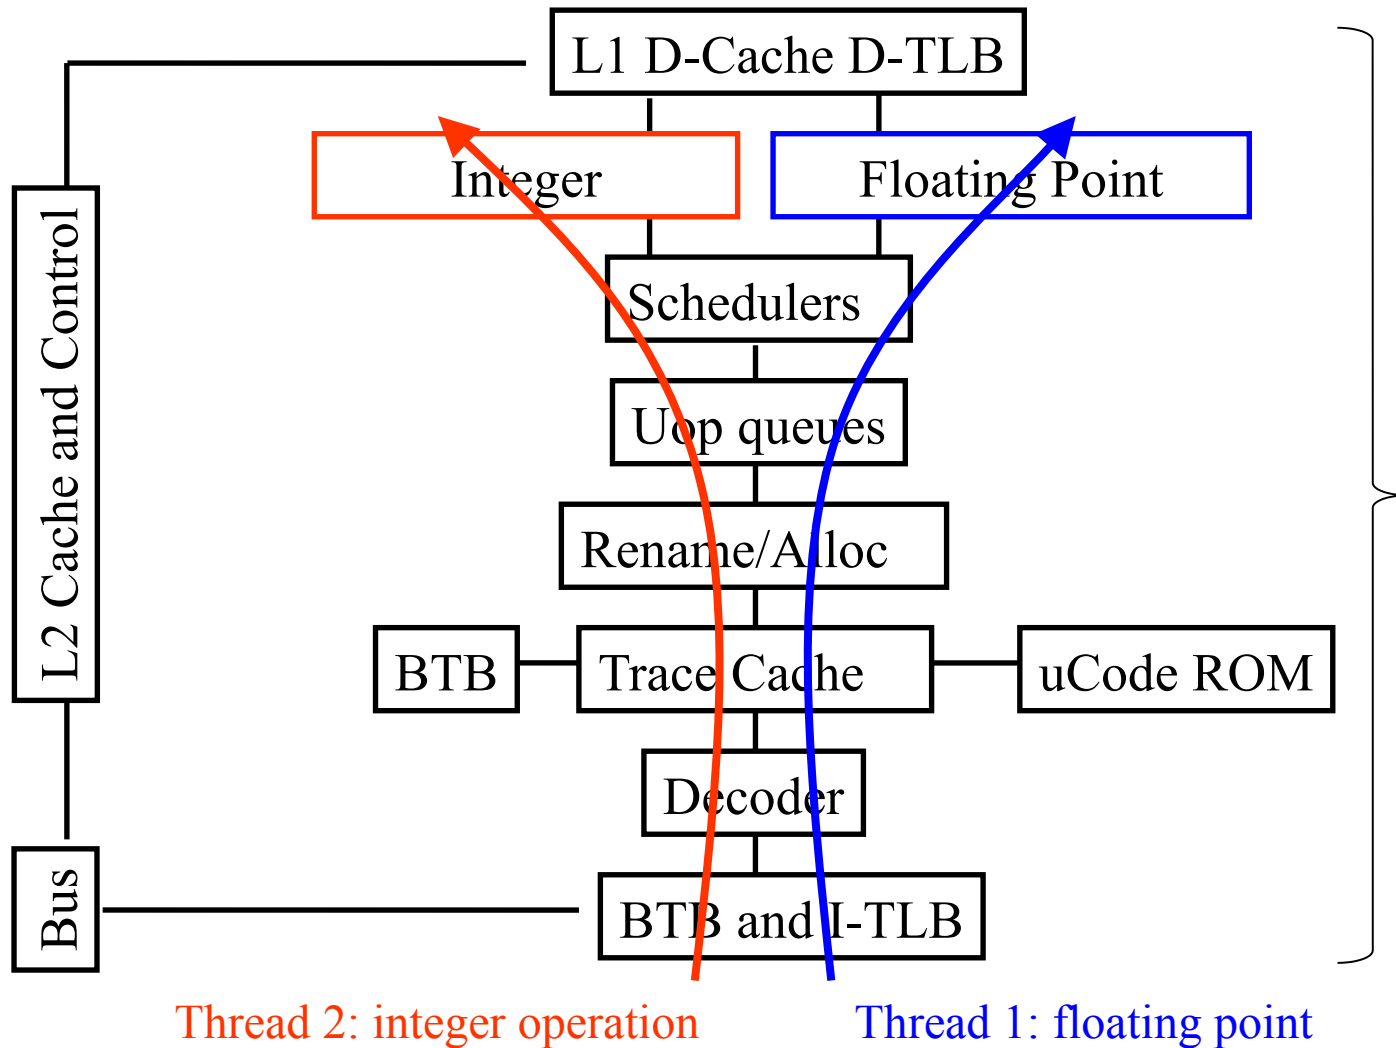
Hiperthreading



Hiperthreading

- Processing can be stalled...
 - Wait from fp result
 - Wait from memory (cache miss)
- In SMT–Simultaneous Multithreading, several thread execute concurrently in the same core, but:
 - On thread processes integred other floats
- Bubles on the piple from one thread
 - Are used by the other threads
- 2 “Virtual” cores per real core
 - /proc/cpuinfo (processor, cpu cores, core id)

Hipertreading



Hyperthreading

- First intel Simultaneous MT
 - 2002 Xeon, the Pentium 4
 - Hyper-Threading.
- Gains of 15%–30%
 - 5% increase of CPU area
- Each core contains
 - 2 Logical processor
 - Registers , L1 cache of 16KB, Interrupt control.
- 1 real processor
 - System BUS, L2 cache, ALU, FPU.