

Synchronization

Access to the same variable

- Multiprogramming in OS allows different processes or threads to access common data
 - Read
 - write
- The order by which tasks are scheduled can produce multiple results
 - Multiple accesses to the shared variable
 - Multiple accesses to the shared data structure
 - Multiple accesses to a device
 - Solved by the kernel :)
 - Multiple accesses to the same file
 - Depends on the granularity

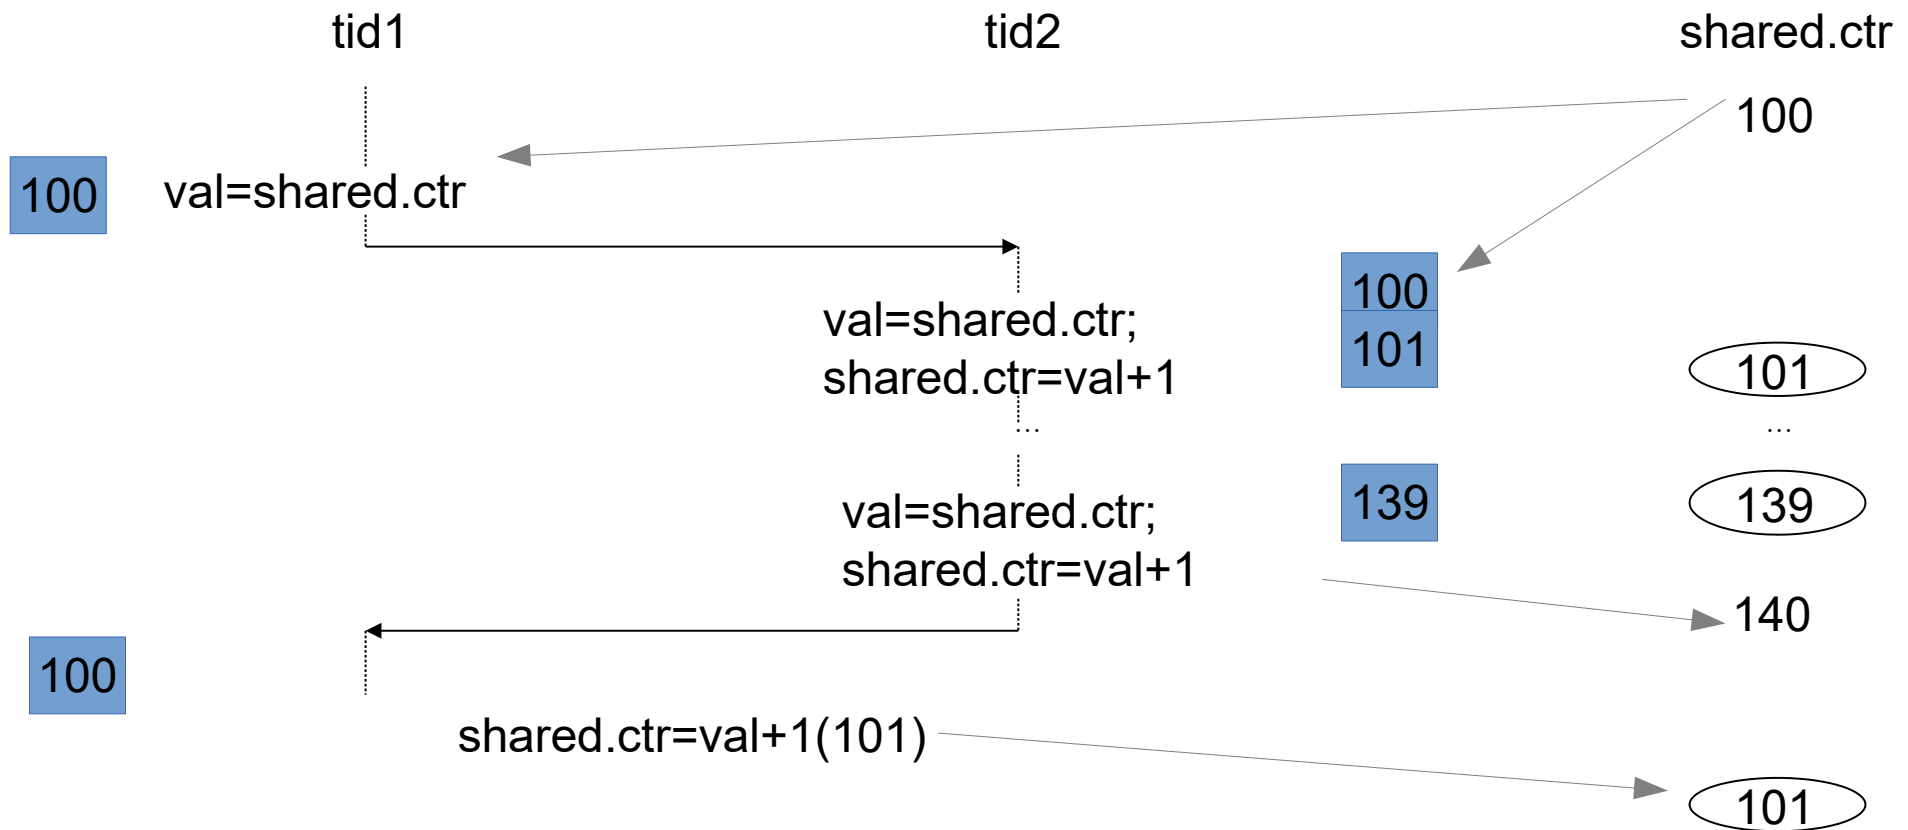
Access to the same variable

```
int shared;

void *count(void *arg) {
    int i, val;
    long this= (long)pthread_self();
    for (i=0; i<NITERS; i++) {
        val=shared;
        printf("%lu: %d\n", this, val);
        shared=val+1;
    }
    return NULL;
}
```

Race condition

- The execution of a thread can be interrupted
 - between
 - `val=shared.ctr;` and `shared.ctr=val+1;`



Access to the same variable

```
int shared;

void *count(void *arg) {
    int i,val;
    long this= (long)pthread_self();
    for (i=0; i<NITERS; i++) {

        printf("%lu: %d\n", this,
               shared ++);

    }
    return NULL;
}
```

Access to the same variable

```
int shared;

void *count(void *arg) {
    int i,val;
    long this= (long)pthread_self();
    for (i=0; i<NITERS; i++) {
        Shared++;

    }
    printf("%lu: %d\n", this,shared);
    return NULL;
}
```

Access to the same variable

- `int shared;`
- - `void *count(void *arg) {`
 - `int i, val;`
 - `long this= (long)pthread_self();`
 - `for (i=0; i<NITERS; i++) {`
 - `printf("%lu: %d\n", this, val);`
 - `shared ++;`
 - `}`
 - `return NULL;`
 - `}`

Access to the same record


- Processes communicate using shared files
 - Two account management functions.

```
int deposit(int account, int amount) {  
    balance = readBalance(account);  
    balance += amount;  
    updateBalance(account, balance);  
    return balance;  
}
```

```
int withdraw(int account, int amount) {  
    balance = readBalance(account);  
    balance -= amount;  
    updateBalance(account, balance);  
    return balance;  
}
```

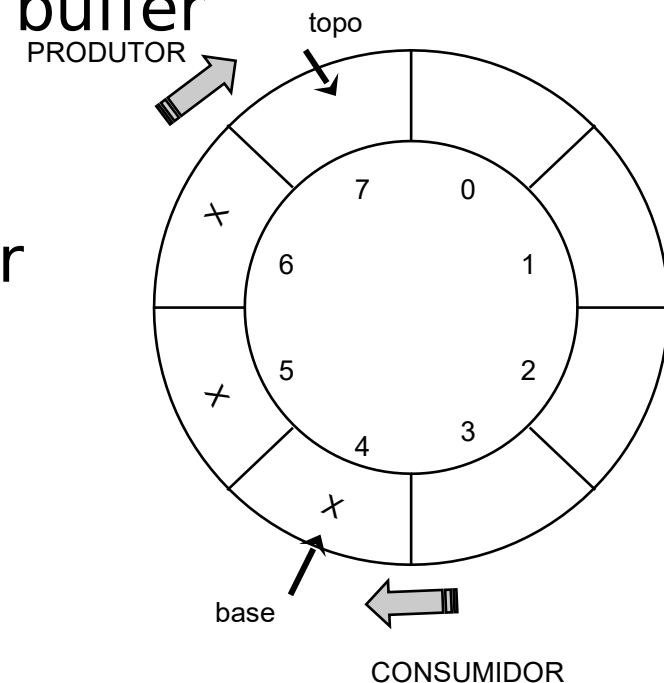

Access to the same variable

- Two concurrent withdrawals
 - Initial value 500€,
 - Withdrawal of 100€



Access to the shared data-structure

- Producer-consumer also present race conditions
 - Two tasks share a common buffer.
 - Producer and consumer
 - The producer write data into the buffer
 - If not full
 - Blocks when buffer is full
 - Consumer reads data from buffer
 - If not empty
 - Blocks if no data is available



Access to the shared data-structure

Producer

```
while(1) {  
    x = create_data()  
    while(counter==SIZE){ }  
    /*buffer cheio*/  
    buffer[topo]=X;  
    counter++;  
    topo=(topo+1)%SIZE;  
}
```

Consumer

```
While(1) {  
    while(counter==0){ }  
    /*buffer vazio*/  
  
    Y=buffer[base];  
    counter--;  
    base=(base+1)%SIZE;  
    consume_data(y)  
}
```

1 producer + 1 consumer

Producer

```
while(1) {  
    x = create_data()  
    while(counter==SIZE){ }  
    /*buffer cheio*/  
    buffer[topo]=X;  
    counter++;  
    topo=(topo+1)%SIZE;  
}
```

MOV %EAX,counter

INC %EAX

MOV counter,%EAX

Consumer

```
While(1) {  
    while(counter==0){ }  
    /*buffer vazio*/  
  
    Y=buffer[base];  
    counter--;  
    base=(base+1)%SIZE;  
    consume_data(y)  
}
```

MOV %EAX,counter

DEC %EAX

MOV counter,%EAX\

1 producer + 1 consumer

- Example of race conditions
 - with counter==4

Entidade	INstrução	Registo	Counter
produtor	MOV %EAX,counter	EAX ← 4	4
produtor	INC %EAX	EAX ← 5	4
consumidor	MOV %EAX,counter	EAX ← 4	4
consumidor	DEC %EAX	EAX ← 3	4
produtor	MOV counter, %EAX	EAX ← 5	5
consumidor	MOV counter, %EAX	EAX ← 3	3

1 producer + 1 consumer

- Producer

```
- while(1) {  
-   /* gera dado X */  
-   while(counter==SIZE)  
-       /*buffer cheio*/ ;  
-   buffer[topo]=X;  
-   counter++;  
-   topo=(topo+1)%SIZE;  
- }
```

- Consumer

```
- While(1) {  
-   /* consome dados */  
-   while(counter==0)  
-       /*buffer vazio*/ ;  
-   Y=buffer[base];  
-  
-   counter--;  
-   base=(base+1)%SIZE;  
- }
```

1 producer + 1 consumer

- Solution to race condition
 - with counter==4
- Similar solution solves multiple consumers

Entidade	INstrução	Registo	Counter
produtor	MOV %EAX,counter	EAX ← 4	4
produtor	INC %EAX	EAX ← 5	4
produtor	MOV counter, %EAX	EAX ← 5	5
consumidor	MOV %EAX,counter	EAX ← 4	5
consumidor	DEC %EAX	EAX ← 3	5
consumidor	MOV counter, %EAX	EAX ← 3	4

1 producer + 1 producer

- Producer

```
- while(1) {  
-   /* gera dado X */  
-   while(counter==SIZE)  
-       /*buffer cheio*/ ;  
-   buffer[topo]=X;  
-   counter++;  
-   topo=(topo+1)%SIZE;  
- }
```

- Producer

```
- while(1) {  
-   /* gera dado X */  
-   while(counter==SIZE)  
-       /*buffer cheio*/ ;  
-  
-   buffer[topo]=X;  
-   counter++;  
-   topo=(topo+1)%SIZE;  
- }
```


1 producer + 2 consumer

- Producer

```
- while(1) {  
-  
- while(counter==SIZE)  
-   /*buffer cheio*/ ;  
-   buffer[topo]=X;  
-   counter++;  
-   topo=(topo+1)%SIZE;  
- }
```

- Consumer

```
- While(1) {  
-  
- while(counter==0)  
-   /*buffer vazio*/ ;  
-   Y=buffer[base];  
-   counter--;  
-   base=(base+1)%SIZE;  
- }
```

Producer / consumer

- Active wait uses CPU cycles
 - OS should interrupt task
 -
-
- The consumer should be awoken when
 - Producer inserts value in memory
- The producer should be awoken when
 - Consumer remove a value from memory

Producer-consumidor

- Producer

```
- while(1) {  
-     /* gera dado X */  
-     if (counter==SIZE)  
-         pause() ;  
-     buffer[topo]=X;  
-     counter++;  
-     if (counter==1)  
-         signal(consumer);  
-     topo=(topo+1)%SIZE;  
- }
```

- Consumer

```
- While(1) {  
-  
-     if (counter==0)  
-         pause() ;  
-     Y=buffer[base];  
-  
-     counter--;  
-     if (counter==N-1)  
-         signal(producer);  
-     base=(base+1)%SIZE;  
- }
```

Pause signal considerations

- The pause and signal actions also show race conditions
- Example:
 - The consumer reads the last value in memory
 - Since `counter == 0`
 - The consumer is paused by the OS (before the `pause()`)
 - Producer starts working
 - Producer inserts an element
 - Since `counter == 1`,
 - executes **`signal(consumer)`**
 - That has no effect since the consumer has not executed the **`pause()`**.
 - Consumer resumes executions
 - Gets blocked in the **`pause()`**
 - Producer continues producing
 - Filling the buffer
 - Not notifying the consumer

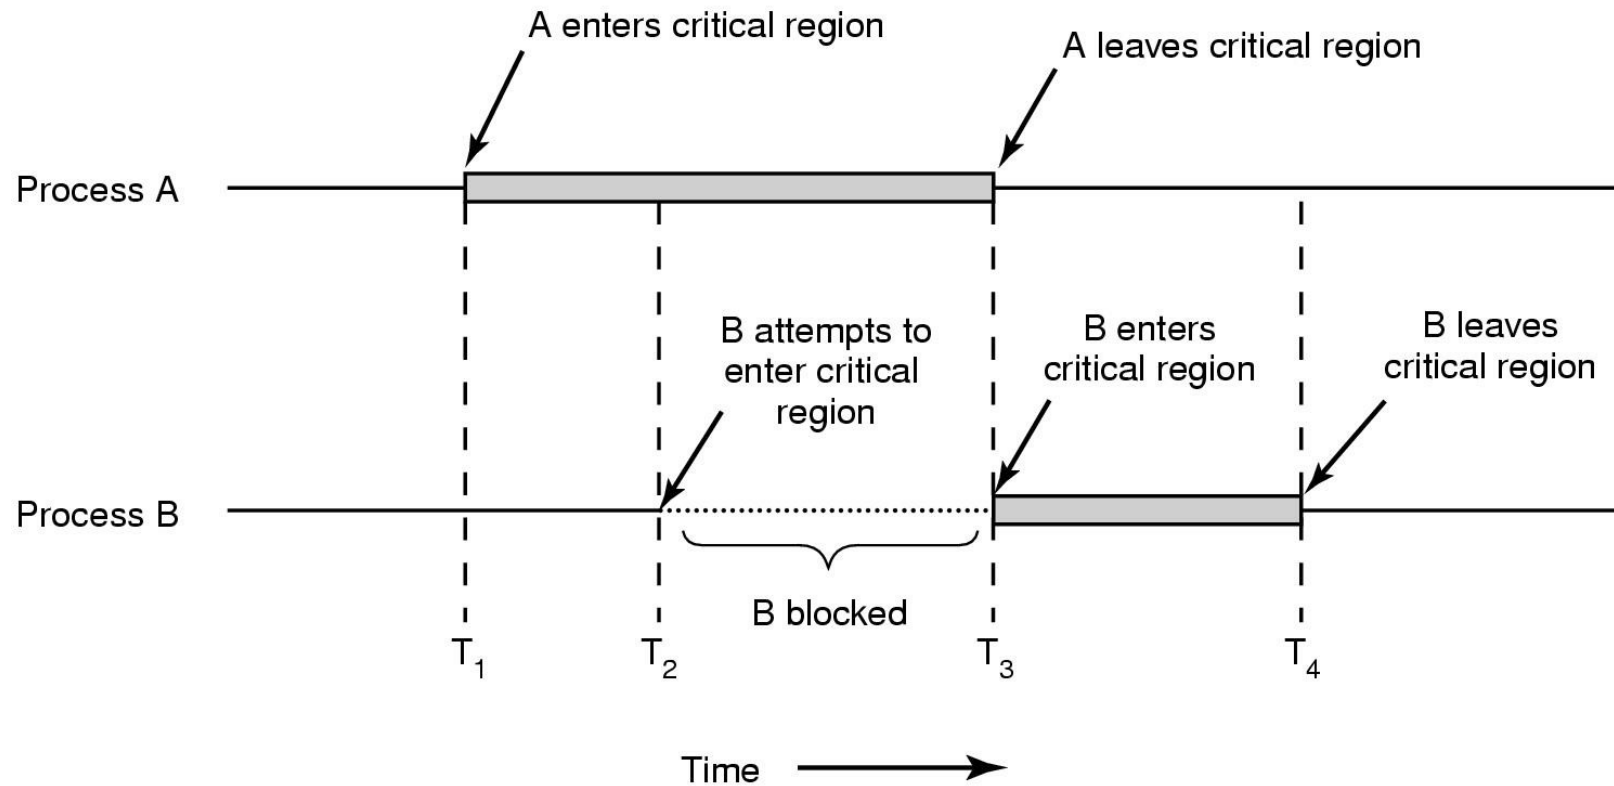
Race condition

- “race condition”
 - Occurs when the result of an execution on shared data
 - depends on the order of the interleaving of instructions.
- A race occurs when two threads can access (read or write) a data variable simultaneously and at least one of the two accesses is a write.
 - (Henzinger 04)
- Results from the fact that the result depends on the last task to conclude.
- To avoid race conditions,
 - Tasks should be synchronized
- Synchronization forces the order of events
 - Special functions
 - communication

Critical region

- Piece of code where resources are shared
 - That can only be executed by one tasks at a time
- Is delimited by read/write instructions to the shared resource
-
- Concurrent reads do not produce inconsistent/incoherent results
- If a task is inside a critical region
 - Other task trying to enter should be blocked

Critical region



Mutual exclusion

- Mechanism that assures that only one task (at most) is inside a critical region
 - No more than one task is executing the critical region code
 - All other tasks are
 - Running non critical code
 - blocked
- Consequences of mutual exclusion
 - starvation (live-lock)
 - A task is able to be executed, but is never scheduled
 - deadlock
 - Due to coding problems, several tasks are waiting for being unblocked by other tasks that are in the same state

Mutual exclusion

- Existing solutions
 - Active wait algorithms (Peterson, Lamport)
 - Hardware
 - Processor special instructions
 - Synchronization objects
 - semaphores, mutexes
 - Using OS services
 - monitors, message passing

Critical regions

- Requirements to be satisfied
 - Mutual exclusion
 - Only one task can be inside the Critical region
 - Progress
 - A task inside the Critical region can not block other task from enetring.
 - Limited wait
 - A task should wait a limited amount of time before entering
- A task should remain inside the Critical Region for a limited time
- The speed and number of processor is undefined
 - Should work for any combination

Critical region programming

```
do {  
    /* secure zone */  
    EnterRegion()  
    Critical region  
    Process(shared_data)  
    LeaveRegion()  
    Remaining code  
    /* secure zone */  
} while(X);\
```

- Entry Region

- Takes request permission to enter
- Task get blocked

- Out region

- Task signal exit of region
- To allow other tasks to enter

- Remaining code

- Code that does not access shared data

Mutual exclusion?

- `void *count(void *arg) {`
- `int i, val;`
- `long this= (long)pthread_self();`
- `for (i=0; i<NITERS; i++) {`
- `printf("%lu: %d\n", this, val);`
- `shared ++;`
- `}`
- `return NULL;`
- `}`

Mutual exclusion?

- Processes communicate using shared files
 - Two account management functions.

```
int deposit(int account, int amount) {  
    balance = readBalance(account);  
    balance += amount;  
    updateBalance(account, balance);  
    return balance;  
}
```

```
int withdraw(int account, int amount) {  
    balance = readBalance(account);  
    balance -= amount;  
    updateBalance(account, balance);  
    return balance;  
}
```

Mutual exclusion?

- Producer

```
- while(1) {  
-   /* gera dado X */  
-   while(counter==SIZE)  
-       /*buffer cheio*/ ;  
-   buffer[topo]=X;  
-   counter++;  
-   topo=(topo+1)%SIZE;  
- }
```

- Consumer

```
- While(1) {  
-     /* consome dados */  
-     while(counter==0)  
-         /*buffer vazio*/ ;  
-     Y=buffer[base];  
-  
-     counter--;  
-     base=(base+1)%SIZE;  
- }
```

POSIX Synchronization primitives

- Synchronization variables
 - Mutexes ✓
 - Spinlocks
 - Read/write locks
 - Semaphores ✓
 - Conditional variables ✓
- Synchronization mechanics
 - Critical region implementations ✓
 - Using synchronization variable
 - Using messages
 - Rendez-vous ✓
 - Barriers ✓
 - Monitors
 - e.g. Java

Mutexes definition

- variable that can only have two values and a list of waiting tasks
 - locked and unlocked
- In linux
 - Introduced in Kernel 2.6.16, corresponding to binary semaphores (0-locked, 1-unlocked).
- Lighter than semaphores
 - semaphore takes 28 bytes vs 16bytes for mutex
 - faster
- A locked mutex is owned by a single task
 - only the task that locked it can unlock
- tasks that try to lock a mutex are stored in the list
 - order of unlock depended on the scheduling policies
- mutex ::= MUTual EXclusion

Mutex definition

- Mutex usage
 - Creation and initialization
 - several tasks execute Entry region
 - try to lock the mutex
 - only on continues
 - becomes the owner of the mutex
 - the owner executes the critical region
 - the owner enters the exit region and unlocks the mutex
 - Other task waiting in the entry region locks the mutex and becomes owner
 - ...
- Mutex is destroyed

POSIX Mutexes

- POSIX mutexes are associated to pthreads:
 - include file `#include <pthread.h>`
 - data type: `pthread_mutex_t mux;`
- A mutex should be initialized before used
- `int pthread_mutex_init(pthread_mutex_t *restrict mutex,`
 - `const pthread_mutexattr_t *restrict attr);`
 - mutex - pointer to variable
 - attr - mutex attributes
 - `PTHREAD_MUTEX_NORMAL` `PTHREAD_MUTEX_ERRORCHECK`
`PTHREAD_MUTEX_RECURSIVE` `PTHREAD_MUTEX_DEFAULT`
- `mux=PTHREAD_MUTEX_INITIALIZER;`
 - static default creation parameters (attr = NULL)
 - more efficient

POSIX Mutexes

- A mutex is destroyed by:
- **int pthread_mutex_destroy(pthread_mutex_t *);**
 - int error;
 - pthread_mutex_t mux;
 -
 - if (error=**pthread_mutex_init**(&mux,NULL))
 - perror("mutex_init: ")
 - /* ... */
 - if (error=pthread_mutex_destroy(&mux))
 - perror(mutex_destroy: ") ;

POSIX Mutexes

- mutex locking
- **`int pthread_mutex_lock(pthread_mutex_t *mutex);`**
 - block thread until resource is available.
 - returns when task enters critical region
- **`int pthread_mutex_trylock(pthread_mutex_t *mutex);`**
 - retorna imediatamente.
- both functions return 0 in success
- `pthread_mutex_trylock`
 - if thread can not lock
 - error variable equals EBUSY
- Mutexes should be kept locked for the minimum amount of time

POSIX Mutexes

- Mutex unlock
- **`int pthread_mutex_unlock(pthread_mutex_t *mutex);`**
- Critical region can be guaranteed by
 - `pthread_mutex_t mux=PTHREAD_MUTEX_INITIALIZER;`
 - `do{`
 - `pthread_mutex_lock(&mux); /* RE */`
 - `/* RC */`
 - `pthread_mutex_unlock(&mux); /* RS */`
 - `/* RR */`
 - `while(TRUE);`

POSIX Mutexes

- Mutexes attributes
 - pthread_mutexattr_*
- protocol
 - PTHREAD_PRIO_INHERIT
 - thrd1 priority and scheduling are affected when higher-priority threads block on one or more mutexes owned by thrd1
 - PTHREAD_PRIO_NONE
 - A thread's priority and scheduling are not affected by the mutex ownership.
 - PTHREAD_PRIO_PROTECT
 - thrd2 priority and scheduling is affected when the thread owns one or more mutexes
- Shared
 - boolean - allow mutex to be shared among threads from more than one process

POSIX Mutexes

- Mutexes robustness
- PTHREAD_MUTEX_STALLED
 - No special actions are taken if the owner of the mutex is terminated while holding the mutex lock.
 - This is the default value.
- PTHREAD_MUTEX_ROBUST
 - If the process containing the owning thread of a robust mutex terminates while holding the mutex lock,
 - the next thread that acquires the mutex shall be notified about the termination by the return value

POSIX Mutexes

- Mutexes type
- PTHREAD_MUTEX_NORMAL
 - This type of mutex does not detect deadlock.
 - A thread attempting to relock this mutex without first unlocking it will deadlock.
- PTHREAD_MUTEX_ERRORCHECK
 - This type of mutex provides error checking.
 - A thread attempting to relock this mutex without first unlocking it will return with an error.
- PTHREAD_MUTEX_RECURSIVE
 - A thread attempting to relock this mutex without first unlocking it will succeed in locking the mutex.
 - Multiple locks of this mutex require the same number of unlocks to release the mutex before another thread can acquire the mutex.
- PTHREAD_MUTEX_DEFAULT
 - Attempting to recursively lock a mutex of this type results in undefined behaviour.

POSIX Mutexes

Mutex Type	Robustness	Relock	Unlock When Not Owner
NORMAL	non-robust	deadlock	undefined behavior
NORMAL	robust	deadlock	error returned
ERRORCHECK	either	error returned	error returned
RECURSIVE	either	recursive	error returned
DEFAULT	non-robust	undefined behavior	undefined behavior
DEFAULT	robust	undefined behavior	error returned

Spin Locks

Spin Locks

- Spin locks are a low-level synchronization mechanism
 - suitable for shared memory multiprocessors.
- When the calling thread requests a spin lock that is already held by another thread
 - the second thread spins in a loop to test if the lock has become available.
- When the lock is obtained,
 - it should be held only for a short time,
 - as the spinning wastes processor cycles.
- Callers should unlock spin locks before calling sleep operations to enable other threads to obtain the lock.

Spin Locks

- initialization
 - `int pthread_spin_init(pthread_spinlock_t *lock, int pshared);`
- locking
 - `int pthread_spin_lock(pthread_spinlock_t *lock);`
 - `int pthread_spin_trylock(pthread_spinlock_t *lock);`
- unlocking
 - `int pthread_spin_unlock(pthread_spinlock_t *lock);`
- Destroying
 - `int pthread_spin_destroy(pthread_spinlock_t *lock);`

Read-Write Locks

Read-Write Lock

- Read-write locks permit concurrent reads and exclusive writes to a protected shared resource.
- The read-write lock is a single entity that can be locked in read or write mode.
- To modify a resource, a thread must first acquire the exclusive write lock.
 - An exclusive write lock is not permitted until all read locks have been released.
-
- Read-write locks support concurrent reads of data structures
 - read operation does not change the record's information.
- When the data is to be updated
 - the write operation must acquire an exclusive write lock.

Not fair

- Th1 ReadLock (enter)
- Th2 ReadLock (enter)
- Th3 WriteLock (blocked)
- Th2 unlock (leave)
- Th4 ReadLock (enter)
- Th1 unlock (leave)
- Th2 ReadLock (enter)
-
- Th3 is in starvation

- Th1 ReadLock (enter)
- Th2 ReadLock (enter)
- Th3 WriteLock (blocked)
- Th2 unlock (leave)
- Th4 ReadLock (blocked)
- Th1 unlock (leave)
 - Th3 enters
- Th2 ReadLock (blocked)

Read-Write Lock

- Initialization

- `int pthread_rwlock_init(pthread_rwlock_t *restrict rwlock,`
– `const pthread_rwlockattr_t *restrict attr);`
 - `rwlock` will contain the reference to the lock
 - `attr` can be `NULL`

- Destruction

- `int pthread_rwlock_destroy(pthread_rwlock_t *rwlock);`

- locking

- `int pthread_rwlock_rdlock(pthread_rwlock_t *rwlock);`
– `int pthread_rwlock_wrlock(pthread_rwlock_t *rwlock);`

- unlocking

- `int pthread_rwlock_unlock (pthread_rwlock_t *rwlock);`

Semaphores

- Abstraction for a counter and a task list, used for synchronization purposes.
 - Proposed by Dijkstra in 1965.
 - Do not require active wait.
 - All modern OS include a version of semaphores (Unix, Windows, ...)
- `typedef struct{`
- `unsigned counter;`
- `processList *queue;`
- `}sem_t;`

Semaphores

- S.counter
 - Defines how many tasks can pass the semaphore without blocking
- If s.coutner = 0
 - The next task to try to enter will get blocked
- If s.coutner is always ≤ 1
 - Mutual exclusion is guaranteed
 - The length of the queue depends on the number of tasks waiting to enter
-

Semaphores

- `down(S)`
 - Used by a task when trying to access a resource
 - Access is given by another task issuing `up(S)`.
 - Task can get blocked if capacity is full (`counter == 0`)
- `up(S)`
 - Used by a task to signal the resource availability
 - `up(S)` is not blocking

Semaphores

- Critical region
 - Initialize a semaphore with counter = 1
 - Do a down when entering the critical region
 - Do a Up when leaving the critical region
- Rendezvous
 - Initialize a semaphore with counter = 0
 - The tasks that does down(S) will get blocked
 - Until other task does the UP(S)
 - Two tasks reandez-vous and continued together

Semaphores

- Wait(S), down(S), or P(S)
 - if (S.counter > 0)
 - S.counter--;
 - else
 - Block(S); /* insert S into the queue */
- Signal(S), up(S), or V(S)
 - if (S.queue != NULL)
 - WakeUp(S); /* removes a process from queue */
 - else
 - S.counter++;
- Wakeup and block depend on the OS
- P(S) e V(S) come from dutch words prolaag (decrement) and verhoog (incrementar)

POSIX Semaphores

- A semaphore is an integer whose value is never allowed to fall below zero.
- POSIX includes a set of functions
 - Man `sem_overview`
 - `#include <semaphore.h>`
 - Link program with `-lpthread`

POSIX Semaphores

- Two operations can be performed on semaphores:
 - increment the semaphore value by one (`sem_post(3)`);
 - and decrement the semaphore value by one (`sem_wait(3)`).
- If the value of a semaphore is currently zero, then a `sem_wait(3)` operation will block until the value becomes greater than zero.

POSIX Semaphores

- POSIX offer two forms of semaphores with respect to creation
 - Named semaphores
 - Identified by a global name (null terminated string started with /
 - Multiple processes can operate on the same named semaphore by passing the same name to `sem_open`
 - Unnamed (memory-based semaphores)
 - Created in memory shared by multiple threads or processes

POSIX Semaphores

- Posix offers two sharing mechanism
 - Not shared
 - Just threads of the process
 - Shared
 - Unnamed - Shared by threads in related processes parent children or using shared memory
 - Named - Shared by threads in multiple processes
 - Semaphores are also classified by the maximum number of tasks (N) that can access the resources
 - If N equals 1 => binary semaphore

POSIX Semaphores

Named		unnamed
sem_open()		sem_init()
	sem_wait() sem_trywait() sem_post() sem_getvalue()	
sem_close() sem_unlink()		sem_destroy()

Service	POSIX
up	sem_post
down	sem_wait

POSIX unnamed Semaphores

- An unnamed semaphore is a variable of type **sem_t**
 - #include <semaphore.h>
 - sem_t sem;
- Should be initialized before being used by
 - Child processes,
 - threads.
- **int sem_init(sem_t *sem, int pshared, unsigned int value)**
 - sem_init() initializes the unnamed semaphore at the address pointed to by **sem**.
 - The **pshared** argument indicates whether this semaphore is to be shared between the threads of a process, or between processes.
 - The **value** argument specifies the initial value for the semaphore.

POSIX unnamed Semaphores

- Unused semaphore should be destroyed
 - `int sem_destroy(sem_t *);`
-
- `sem_t semaforo;`
- `if (sem_init(&semaforo, 0, 1) == -1)`
- `perror("Falha na inicializacao");`
-
- `if (sem_destroy(&semaforo) == -1)`
- `perror("Falha na eliminacao");`

POSIX named Semaphores

- Allows the synchronization of processes without shared memory
 - Use a global name
 - string following the form **/name**
- Named semaphores are installed/located in `/dev/shm`, with the name **sem.name**
- A semaphore should be opened before used

POSIX named Semaphores

- `sem_t *sem_open(const char *name, int oflag)`
- `sem_t *sem_open(const char *name, int oflag, mode_t mode, unsigned int value)`
 - `sem_open()` creates a new POSIX semaphore or opens an existing semaphore.
 - The semaphore is identified by name.
 - The `oflag` argument specifies flags that control the operation of the call.
 - If `O_CREAT` is specified in `oflag`, then the semaphore is created if it does not already exist.
- If `O_CREAT` is specified in **`oflag`**, two additional arguments must be supplied
 - **`mode_t`**, permissions (owner, group, user)
 - The value argument specifies the initial value for the new semaphore..
- If **`oflag`** contains `O_CREAT` and `O_EXCL`
 - Error if semaphore already exists
- If **`oflag`** is `O_CREAT` and semaphore exists
 - 3rd and 4th parameters are ignored

POSIX named Semaphores

- When the semaphore is of no use should be closed
 - `int sem_close(sem_t *);`
- The last process should
 - Close the semaphore (`sem_close`)
 - Remove the corresponding file
 - `int sem_unlink(const char *);`
- If a process maintains the semaphore opened
 - `sem_unlink` is blocked until the semaphore is closed
- If the semaphore is not closed/unlinked
 - New uses of the semaphore are undefined....

POSIX Semaphores

- down(S) is implemented by the function
 - `int sem_wait(sem_t *);`
 - If counter is zero
 - The thread executing the function is blocked.
 - Remaining thread in the process continue executing
- up(S) is implemented by
 - `int sem_post(sem_t *);`

Mutual exclusion

- `sem_init(&sem, 0, 1)`
 - Or
- `Sem = sem_open(... , O_CREAT, ... , 1);`
- `do {`
- `sem_wait(&sem); /* RE */`
- `/* RC */`
- `sem_post(&sem); /* RS */`
- `/* RR */`
- `} while (TRUE);`

Semaphores

- A thread can pool the value of a semaphore
 - `int sem_getvalue(sem_t *sem, int *sval);`
 - `sem_getvalue()` places the current value of the semaphore pointed to `sem` into the integer pointed to by `sval`.
- If one or more processes or threads are blocked waiting to lock the semaphore with `sem_wait(3)`,
 - POSIX.1-2001 permits two possibilities for the value returned in `sval`:
 - either 0 is returned;
 - negative number whose absolute value is the count of the number of processes and threads currently blocked in `sem_wait(3)`.
 - Linux adopts the former behavior (return 0)

POSIX Semaphores

- `int sem_trywait(sem_t *sem);`
 - `sem_trywait()` is the same as `sem_wait()`, except that if the decrement cannot be immediately performed, then call returns an error (errno set to EAGAIN) instead of blocking.
- `int sem_timedwait(sem_t *sem, const struct timespec *abs_timeout);`
 - same as `sem_wait()`, except that **abs_timeout** specifies a limit on the amount of time that the call should block
 - If the timeout has already expired by the time of the call, then `sem_timedwait()` fails with a timeout error (errno set to ETIMEDOUT).

Condition Variables

- We need additional mechanisms to wait inside locked regions
- However, holding the lock while waiting prevents other threads from entering the locked region
- Condition variables make it possible to sleep inside a critical section
 - Atomically release the lock & go to sleep

Condition Variables

- Each condition variable
 - Consists of a queue of threads
 - Provides three operations
 - Wait();
 - Atomically release the lock and go to sleep
 - Reacquire lock on return
 - Signal();
 - Wake up one waiting thread, if any
 - Broadcast();
 - Wake up all waiting threads
- The three operations can only be used inside locked regions

An Example of Using Condition Variables

```
AddToQueue() {  
    lock.Acquire();  
    // put 1 item to the queue  
    condition.Signal(&lock);  
    lock.Release();  
}
```

```
RemoveFromQueue() {  
    lock.Acquire();  
    while nothing on queue  
        condition.Wait(&lock);  
    lock.Release();  
    return item;  
}
```

Condition variables in pthread

- Waiting and signaling on condition variables
 - Data type: `pthread_cond_t data_cond = PTHREAD_COND_INITIALIZER;`
- Routines
 - `pthread_cond_wait(condition, mutex)`
 - Blocks the thread until the specific condition is signalled.
 - Should be called with mutex locked
 - Automatically release the mutex lock while it waits
 - When return (condition is signaled), mutex is locked again
 - `pthread_cond_signal(condition)`
 - Wake up a thread waiting on the condition variable.
 - Called after mutex is locked, and must unlock mutex after
 - `pthread_cond_broadcast(condition)`
 - Used when multiple threads blocked in the condition
 -
 -
-

Condition Variables

- While mutexes implement synchronization by controlling thread access to data, condition variables allow threads to synchronize based upon the actual value of data.
- Without condition variables, The programmer would need to have threads continually polling (usually in a critical section), to check if the condition is met.
- A condition variable is a way to achieve the same goal without polling (a.k.a. “busy wait”)

Condition Variables

- Useful when a thread needs to wait for a certain condition to be true.
- In pthreads, there are four relevant procedures involving condition variables:
 - `pthread_cond_init(pthread_cond_t *cv, NULL);`
 - `pthread_cond_destroy(pthread_cond_t *cv);`
 - `pthread_cond_wait(pthread_cond_t *cv, pthread_mutex_t *lock);`
 - `pthread_cond_signal(pthread_cond_t *cv);`

Creating and Destroying Conditional Variables

- Condition variables must be declared with type `pthread_cond_t`, and must be initialized before they can be used.
 - Statically, when it is declared. For example:
 - `pthread_cond_t myconvar = PTHREAD_COND_INITIALIZER;`
 - Dynamically
 - `pthread_cond_init(cond, attr);`
- `pthread_cond_destroy(cond)`
 - used to free a condition variable that is no longer needed.

pthread_cond_init

- A condition variable is a pthread_cond_t

```
#include <pthread.h>
```

```
pthread_cond_t cond;
```

- Should be initialized before being used

```
int pthread_cond_init(pthread_cond_t *, const  
pthread_condattr_t *);
```

- 1º parâmetro: address of condition variable
 - 2º parâmetro: attributes (can be NULL)
- static initialization
 - cont=PTHREAD_COND_INITIALIZER;

pthread_cond_destroy

- termination

```
int pthread_cond_destroy( pthread_cond_t *);
```

- Condition variable usage:
 - Created and initialized
 - If condition is not satisfied block/ otherwise continue.
 - Other thread changes the condition and
 - signals other threads about the new condition values
 - one by one or broadcast
 - Signaled thread continues execution.
 -
 - Condition variable is terminated.

pthread_cond_wait

- `pthread_cond_wait(cv, lock)`
 - is called by a thread when it wants to block and wait for a condition to be true.
- It is assumed that the thread has locked the mutex indicated by the second parameter.
- The thread releases the mutex, and blocks until awakened by a `pthread_cond_signal()` call from another thread.
- When it is awakened,
 - it waits until it can acquire the mutex,
 - once acquired, it returns from the `pthread_cond_wait()` call.

pthread_cond_wait

- Waiting on a condition variable:

```
int pthread_cond_wait(pthread_cond_t *,  
pthread_mutex_t *);
```

- wait until other thread signals/broadcasts
-

```
int pthread_cond_timedwait(pthread_cond_t *,  
pthread_mutex_t *,const struct timespec*);
```

- Timed wait.

Condition variables

- **pthread_cond_wait**
- **pthread_cond_timedwait**
 - 1st parameter: condition variable
 - 2nd parameter: mutex that guards the critical region.
 - pthread_cond_timedwait
 - 3^o parâmetro: função de temporização da espera
- pthread_cond_wait() e pthread_cond_timedwait()
 - should be called after a thread_mutex_lock().

pthread_cond_signal

- `pthread_cond_signal()`
 - checks to see if there are any threads waiting on the specified condition variable.
 - If not, then it simply returns.
- If there are threads waiting,
 - then one is awakened.
- There can be no assumption about the order in which threads
 - It is natural to assume that they will be awakened in the order in which they waited, but that may not be the case...
- Use loop or `pthread_cond_broadcast()` to awake all waiting threads.


Condition variables

- **pthread_timedwait**
- **pthread_cond_timedwait**
- `pthread_mutex_lock();`

- `while(condition_is_false)`

- `pthread_cond_wait();`

- `pthread_mutex_unlock();`



pthread_cond_signal pthread_cond_broadcast

- the thread can change the condition variable

int pthread_cond_signal(pthread_cond_t *);

- unlocks at least one thread

int pthread_cond_broadcast(pthread_cond_t *);

- unlocks all threads

- Other Thread only resumes after this thread releases mutual exclusion

Condition variables

- `void *consumer(void *) {`
- `while(1) {`
- `pthread_mutex_lock(&mux);`
- `if (var==threshlold) break;`
- `pthread_mutex_unlock(&mux);`
- `}`
- `pthread_mutex_lock(&data_mutex);`
- `<Extract data from queue;>`
- `if (queue is empty)`
- `data_avail = 0;`
- `pthread_mutex_unlock(&data_mutex);`
- `<Consume Data>`
- `}`
-

- `void *consumer(void *) {`
- `pthread_mutex_lock(&data_mutex);`
- **`while(!data_avail);`**
- **`/* do nothing */`**
-
- `<Extract data from queue;>`
- `if (queue is empty)`
- `data_avail = 0;`
- `pthread_mutex_unlock(&data_mutex);`
- `<Consume Data>`
- `}`
-

Condition variables

```
void *producer(void *) {  
    <Produce data>  
    pthread_mutex_lock(&data_mutex);  
    <Insert data into queue;>  
    data_avail = 1;  
  
    pthread_cond_signal(&data_cond);  
    pthread_mutex_unlock(&data_mutex);  
}
```

```
void *consumer(void *) {  
    pthread_mutex_lock(&data_mutex);  
    while( !data_avail ) {  
        /* sleep on condition variable*/  
        pthread_cond_wait(&data_cond,  
                           &data_mutex);  
    }  
    /* woken up */  
    <Extract data from queue;>  
    if (queue is empty)  
        data_avail = 0;  
    pthread_mutex_unlock(&data_mutex);  
    <Consume Data>  
}
```

- Task Consumer 1

- void *consumer(void *) {



- **pthread_mutex_lock(&data_mutex);**



- while(!data_avail) {



- **pthread_cond_wait(&data_cond,**

-

- **&data_mutex);**

-

- }

-

-



- **pthread_mutex_unlock(&data_mutex);**

- <Consume Data>

-

- }

- Task Consumer 2

- void *consumer(void *) {

-

-

-

-

-



- **pthread_mutex_lock(&data_mutex);**

-

-



- while(!data_avail) {



- /* sleep on condition variable*/



- **pthread_cond_wait(&data_cond,**

- **&data_mutex);**

Condition variables

- task Producer

```
- void *producer(void *) {  
  😊 <Produce data>  
  🌑 pthread_mutex_lock(&data_mutex);  
  -  
  -  
  -  
  😊 <Insert data into queue;>  
  😊 data_avail = 1;  
  😊 pthread_cond_signal(&data_cond);  
  😊 pthread_mutex_unlock(&data_mutex  
    );  
  - }  
  -
```

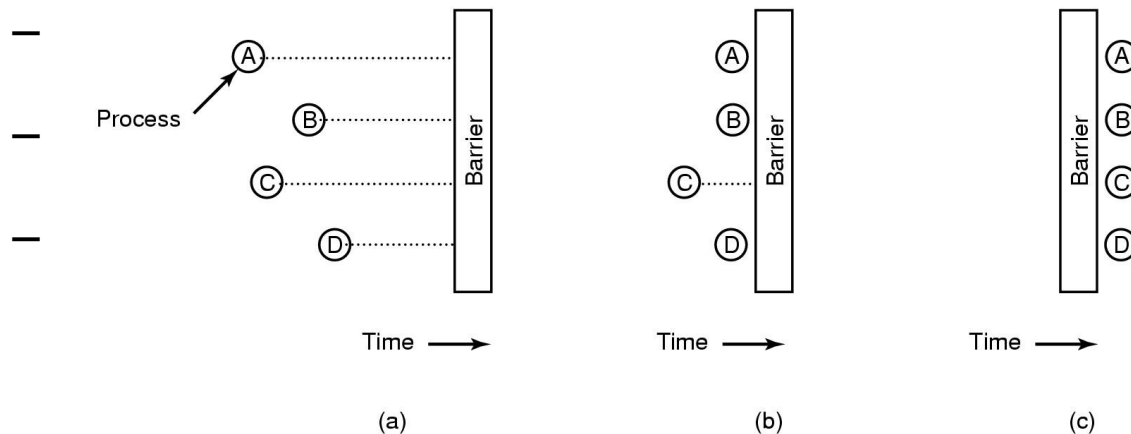
- task Consumer 1

```
- void *consumer(void *) {  
  😊 pthread_mutex_lock(&data_mutex);  
  😊 while( !data_avail ) {  
    -  
    - /* sleep on condition variable*/  
    🌑 pthread_cond_wait(&data_cond,  
      &data_mutex);  
    - }  
    -  
    -  
    😊 if (queue is empty)  
    😊 data_avail = 0;  
    -  
    - pthread_mutex_unlock(&data_mute  
      x);  
    -
```

Barriers

- [Def] Barriers:

- Synchronization mechanism that blocks threads until a defined number of threads arrives at the barrier



- Are used when processing is done in steps, that require the completion of a number of threads
 - The next step is only started (by all threads at the same time) after all threads have finished the preceding step

Barriers

- Pthreads provides the type `pthread_barrier_t`
**`pthread_barrier_init(pthread_barrier_t *,`
`pthread_barrierattr_t *,`
`unsigned int)`**
 - 1st parameter - barrier object
 - 2nd parameter - barrier attributes
 - 3rd parameter - number of threads that must call `pthread_barrier_wait()` before any of them successfully return from the call.
- `pthread_barrier_destroy(pthread_barrier_t *)`**

Barriers

- Wait / Synchronization

```
int pthread_barrier_wait(  
    pthread_barrier_t *)
```

- All threads block
- When the count value (3rd argument from ini) is reached
 - Count thread unblock and start executing
 - Function returns
- Return value
 - one thread - PTHREAD_BARRIER_SERIAL_THREAD,
 - All others - 0

Semaphores using mutex

```
int semaphore_wait (semaphore_t *sem) {  
    sem->count --;  
    if (sem->count < 0)  
        bloquear  
}
```

```
int semaphore_post (semaphore_t *sem) {  
    sem->count ++;  
    if (sem->count <= 0) {  
        Desbloqueia tarefa  
    }  
}
```

Semaphores using mutex

```
int semaphore_wait (semaphore_t *sem) {  
    int res = pthread_mutex_lock(&(sem->mutex));  
    if (res != 0) return res;  
    sem->count --;  
    if (sem->count < 0)  
        pthread_cond_wait(&(sem->cond), &(sem->mutex));  
    pthread_mutex_unlock(&(sem->mutex));  
    return res;  
}
```

Semaphores using mutex

```
int semaphore_wait (semaphore_t *sem) {  
    sem->count --;  
    if (sem->count < 0)  
        bloquear  
}
```

```
int semaphore_post (semaphore_t *sem) {  
    sem->count ;  
    if (sem->count < 0)  
        bloquear  
}
```

Semaphores using mutex

```
int semaphore_post (semaphore_t *sem) {  
    int res = pthread_mutex_lock(&(sem->mutex));  
    if (res != 0) return res;  
    sem->count ++;  
    if (sem->count <= 0) {  
        res = pthread_cond_signal(&(sem->cond));  
    }  
    pthread_mutex_unlock(&(sem->mutex));  
    return res;  
}
```

Bounded buffer

- Implement a queue that has two functions
 - enqueue() - adds one item into the queue. It blocks if queue is full
 - dequeue() - remove one item from the queue. It blocks if queue is empty
- The queue has fixed limit
- How to signal that writes can enqueue?
- How to signal that reads can dequeue?

Bounded buffer

```
enqueue(int val){
    sem_wait(&_fullSem);
    _queue[_tail]=val;  _tail = (_tail+1)%MaxSize;
    sem_post(_emptySem);
}
int dequeue(){
    sem_wait(&_emptySem);
    int val = _queue[_head];  _head = (_head+1)%MaxSize;
    sem_post(_fullSem);
    return val;
}
```

Bounded buffer

- Initialization:

- `sem_init(&_emptySem, 0, 0);`
- `sem_init(&_fullSem, 0, MaxSize);`

```
void enqueue(int val){  
    sem_wait(&_fullSem);  
    mutex_lock(&_amp;mutex);  
    _queue[_tail]=val;  
    _tail = (_tail+1)%MaxSize;  
    mutex_unlock(&_amp;mutex);  
    sem_post(&_amp;emptySem);  
}
```

```
int dequeue(){  
    sem_wait(&_emptySem);  
    mutex_lock(&_amp;mutex);  
    int val = _queue[_head];  
    _head = (_head+1)%MaxSize;  
    mutex_unlock(&_amp;mutex);  
    sem_post(&_amp;fullSem);  
    return val;  
}
```

RW locks

- Multiple readers may read the data structure simultaneously
- Only one writer may modify it and it needs to exclude the readers.
- Interface:
 - ReadLock() – Lock for reading. Wait if there are writers holding the lock
 - ReadUnlock() – Unlock for reading
 - WriteLock() – Lock for writing. Wait if there are readers or writers holding the lock
 - WriteUnlock() – Unlock for writing

RW locks

- How to gurantee just one writer?
 - `void RWLock::writeLock(){`
 - `sem_wait(&_semAccess);`
 - `}`
 - `void RWLock::writeUnlock(){`
 - `sem_post(&_semAccess);`
 - `}`

RW locks

- How to guarantee that read and writers are not at the same time?
 - first reader blocks writer
 - last read unblocks writer

```
void readLock(){
    _nreaders++;
    if( _nreaders == 1 ) {
        //This is the first reader
        //Get sem_Access
        sem_wait(&_semAccess);
        // read
    }
}
```

```
void readUnlock(){
    _nreaders--;
    if( _nreaders == 0 ) {
        //This is the last reader
        //Allow one writer to
        //proceed if any
        sem_post( &_semAccess );
    }
}
```

RW locks

```
void readLock(){
    mutex_Lock( &_mutex );
    _nreaders++;
    if( _nreaders == 1 ) {
        //This is the first reader
        //Get sem_Access
        sem_wait(&_semAccess);
    }
    mutex_unlock( &_mutex );
}
```

```
void readUnlock(){
    mutex_Lock( &_mutex );
    _nreaders--;
    if( _nreaders == 0 ) {
        //This is the last reader
        //Allow one writer to
        //proceed if any
        sem_post( &_semAccess );
    }
    mutex_unlock( &_mutex );
}
```

- Fairness in locking:
 - First-come-first serve
- Mutexes and semaphores are fair.
 - The thread that has been waiting the longest is the first one to wake up.
- Spin locks (active wait) do not guarantee fairness, the one waiting the longest may not be the one getting it
 - This should not be an issue in the situation when one wants to use spin locks, namely low contention, and short lock holding time