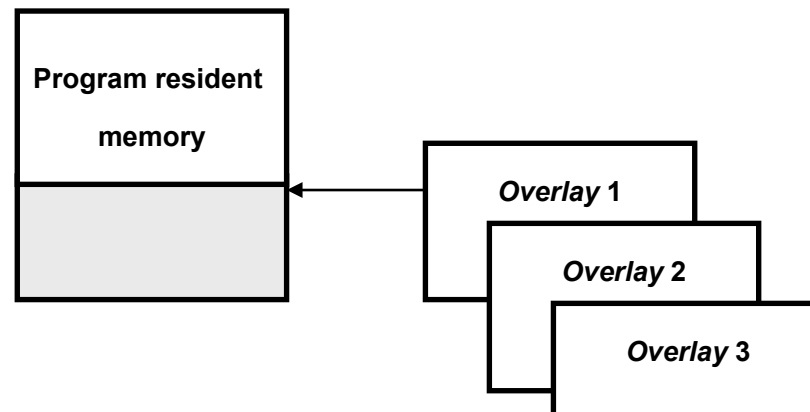


Memory management


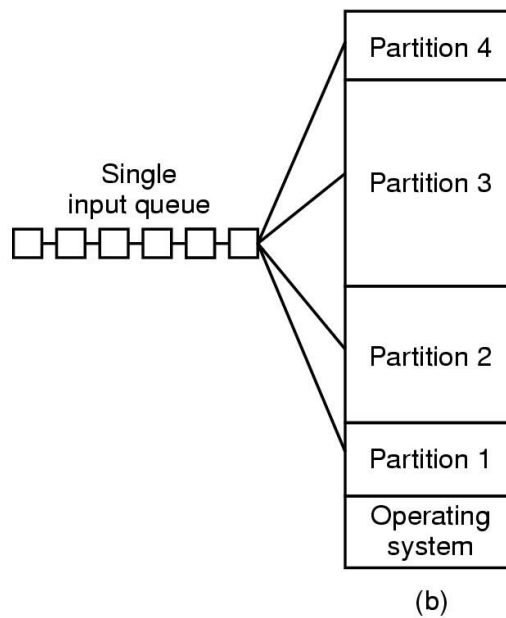
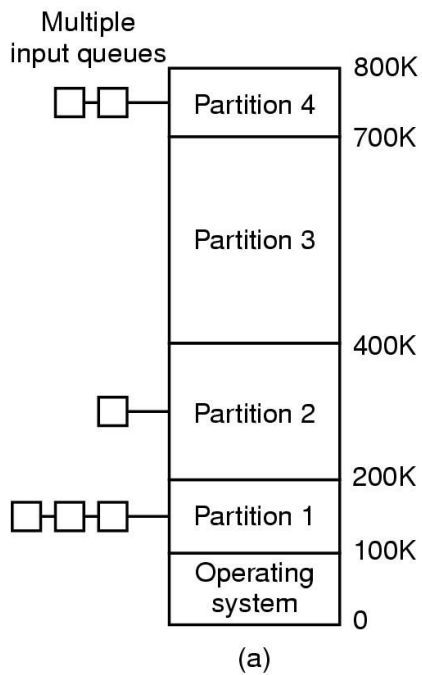
- Single process
 - All memory assigned to the process
 - Addresses defined at compile time
- Multiple processes. How to:
 - assign memory
 - manage addresses?
 - manage relocation?
 - manage program grow?

manage program grow?

- overlay
 - Statically defined memory zone
 - loaded on demand
 - By the program
 - may be replaced



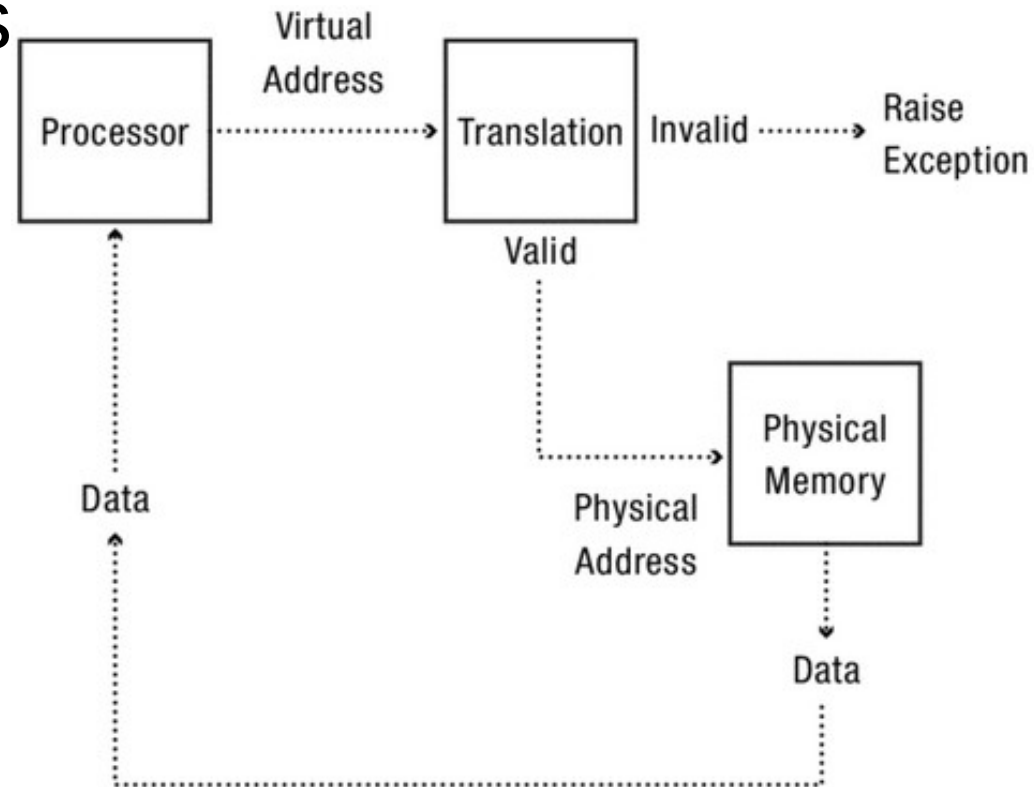
Memory management



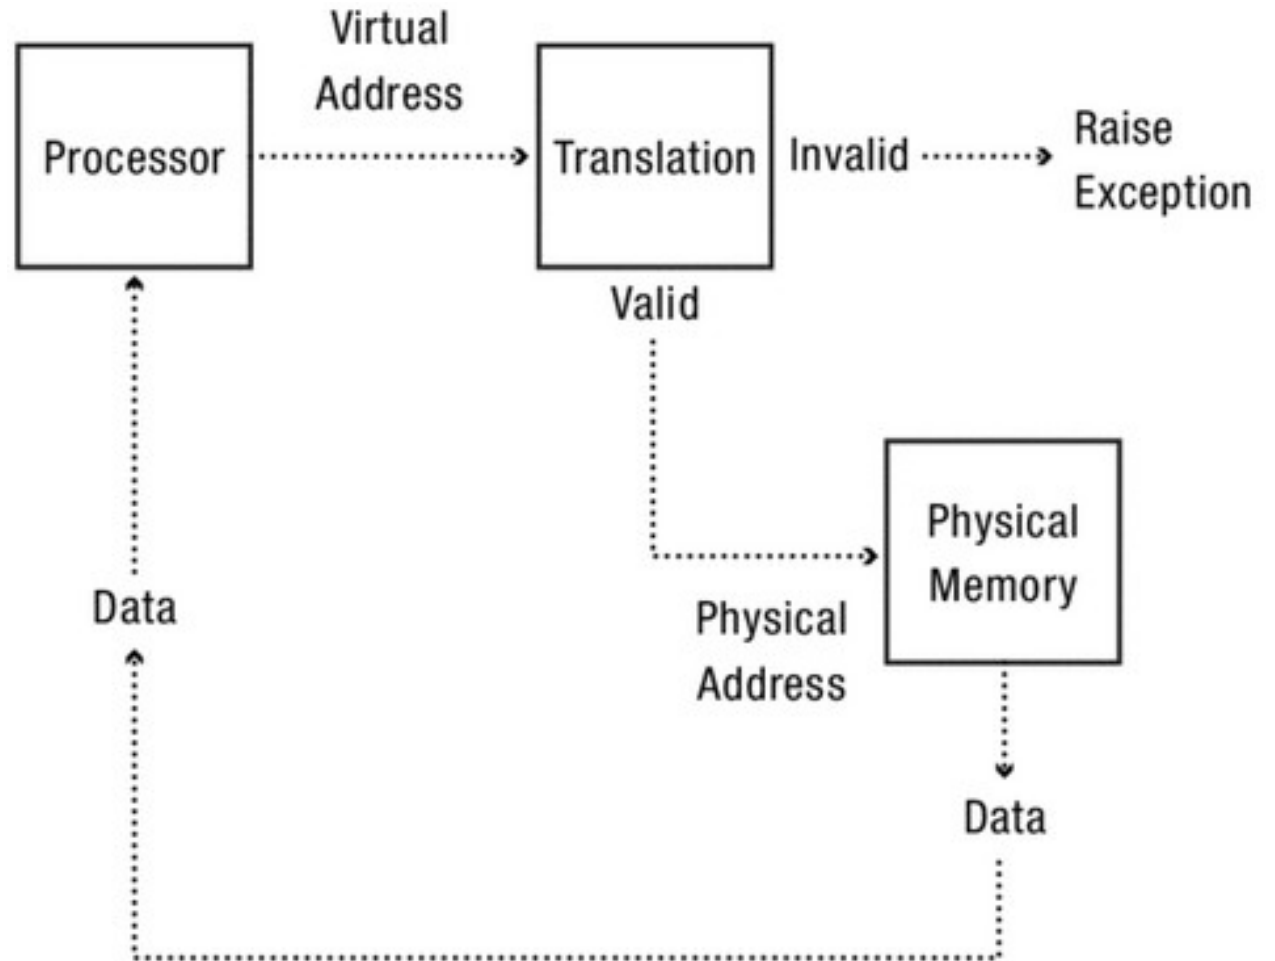
```
for(i = 0; i < 1000; i++){  
    vect[i] = c;  
}  
  
.L3:  
    movl -4(%rbp), %eax  
    cltq  
    movzbl -5(%rbp), %edx  
    movb  %dl, -1008(%rbp,%rax)  
    addl $1, -4(%rbp)  
  
.L2:  
    cmpl $999, -4(%rbp)  
    jle .L3
```

Virtual memory

- Program addresses
 - are independent on the physical location
- memory manager
 - translates addresses
 - virtual -> physical
 - verifies permissions
- Address Translation



Address Translation Concept




Adress translation

- Process isolation
- Interprocess communication
- Shared code
- Program debugging
- Efficient I/O
- Memory mapped files
- Virtual memory
- Checkpointing
- Process migration
- Information flow control
- DSM

Shared memory

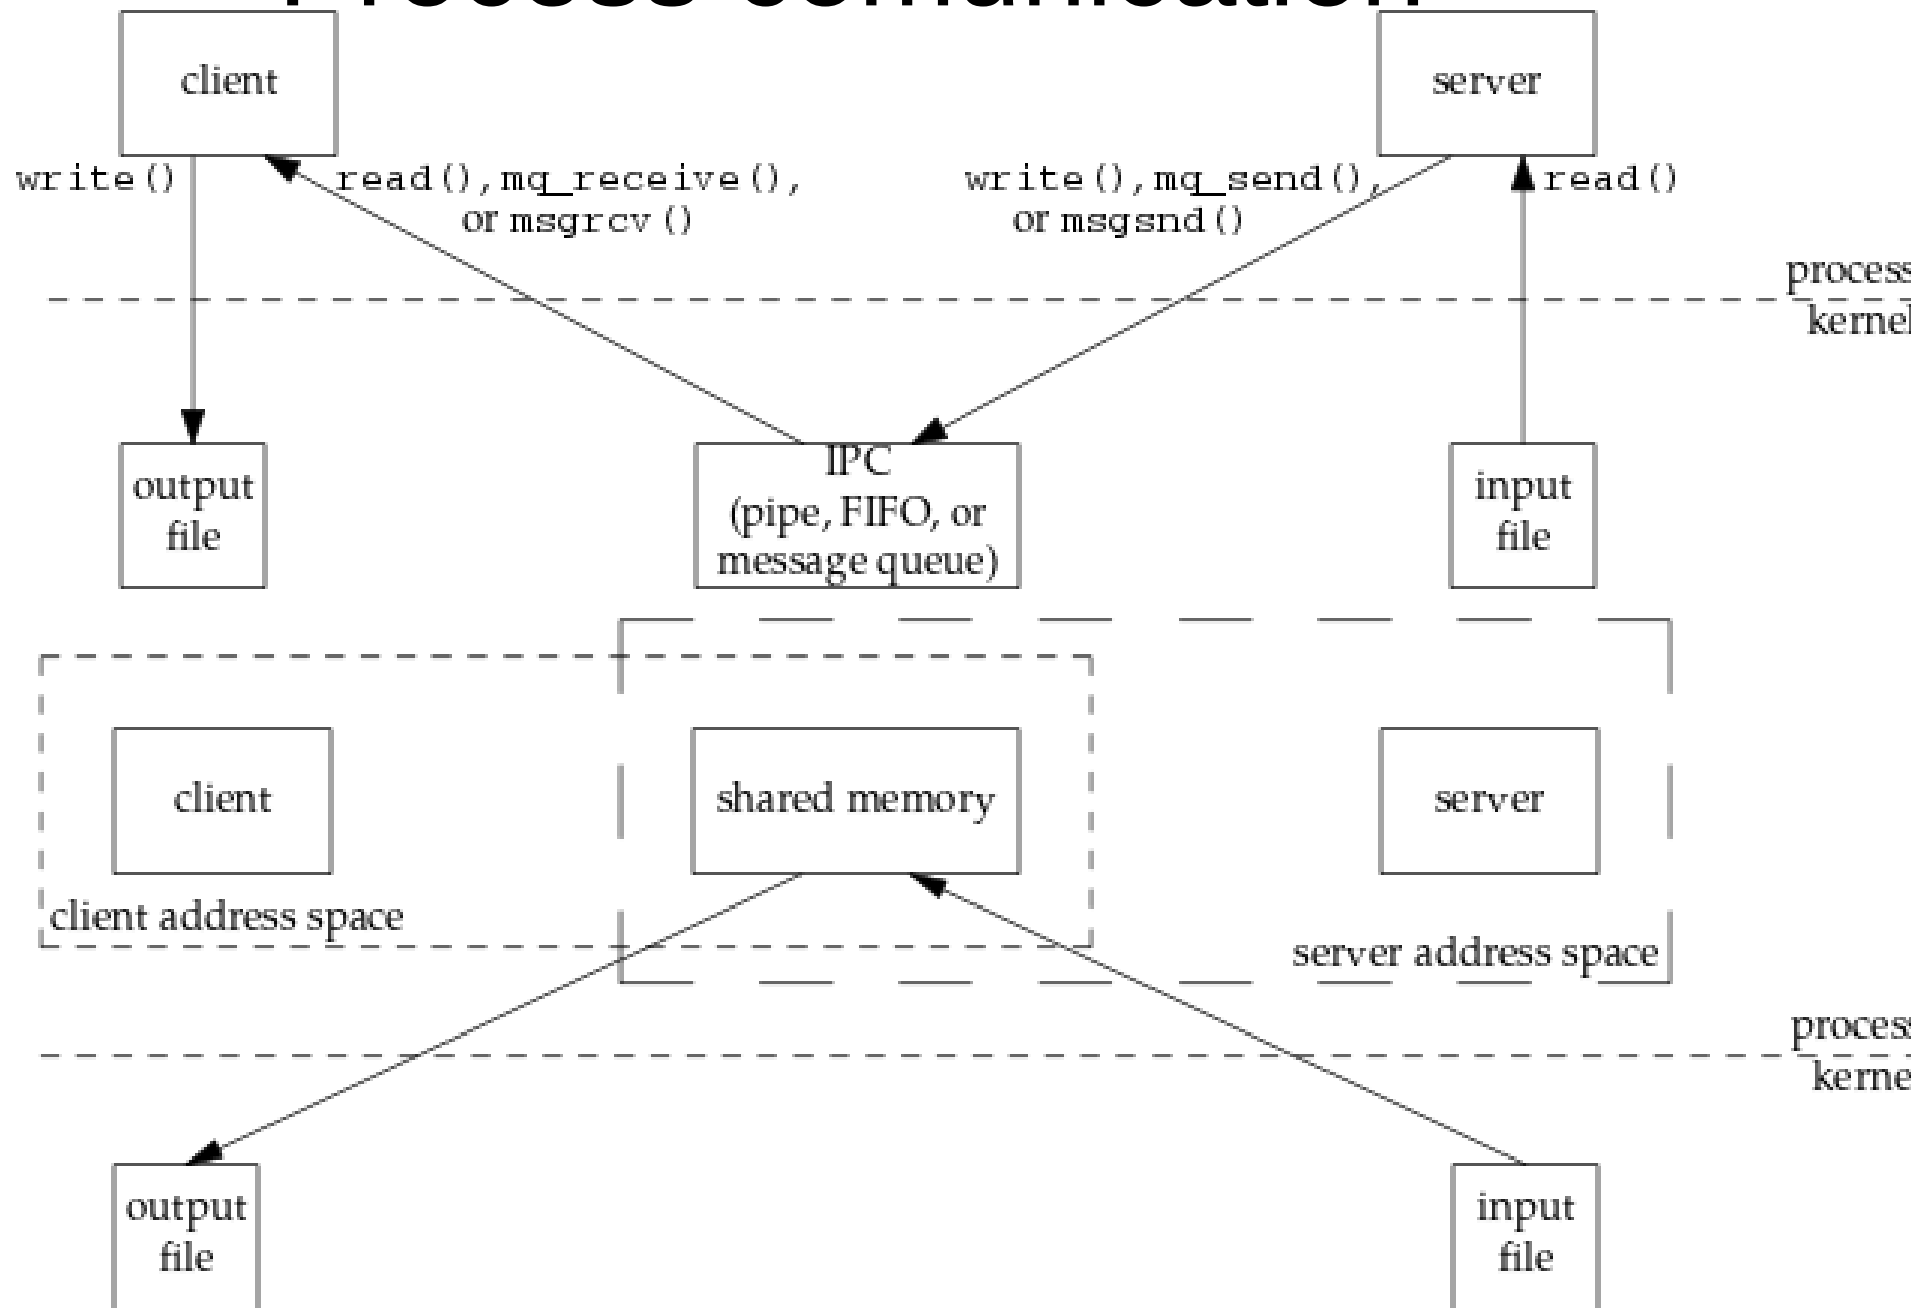
- Processes use regular variables/vector to communicate
 - Available in shared memory
- Should be explicitly created.
 - Each process has its own address space



Shared memory Advantages

- Random Access
 - you can update a small piece in the middle of a data structure, rather than the entire structure
- Efficiency
 - unlike message queues and pipes,
 - copy data between user memory ↔ kernel memory
 - shared memory is directly accessed
 - Shared memory resides in the user process memory
 - Is shared among other processes

Process communication




Shared memory Disadvantages

- No automatic synchronization
 - In pipes or message queues
 - Programmer has to provide synchronization
 - Semaphores or signals.
- Pointers are only valid within a given process.
 - Pointer offsets cannot be assumed to be valid across inter-process boundaries.
 - This complicates the sharing of linked lists or binary trees.
- Variables are “produced by the compiler”
 - Names can not be used to access shared memory

Shared memory in *NIX

- System V shared memory
 - Original shared memory mechanism, still widely used
 - Sharing between unrelated processes
- Shared mappings – mmap
 - Shared file mappings
 - Sharing between unrelated processes, backed by filesystem
 - Shared anonymous mappings
 - Sharing between related processes only (related via fork())
- POSIX shared memory
 - Sharing between unrelated processes, without overhead of filesystem I/O
 - Intended to be simpler and better than older APIs

Shared memory in *NIX

- Programming steps
 - Define Shared data structure
 - Creation of memory segment
 - Configuration
 - Assignment to address
 - Access
 - Disconnection
 - Destruction
- 
- In multiple processes

System V shared memory

- Shared memory operations
 - shmget
 - allocates a shared memory segment
 - shmctl
 - allows the user to receive information on a shared memory segment,
 - set the owner, group, and permissions of a shared memory segment,
 - destroy a segment
 - shmat
 - attaches the shared memory segment (identified by shmid) to the address space of the calling process
 - shmdt
 - detaches the shared memory segment (located at the address specified by shmaddr) from the address space of the calling process

System V shared memory

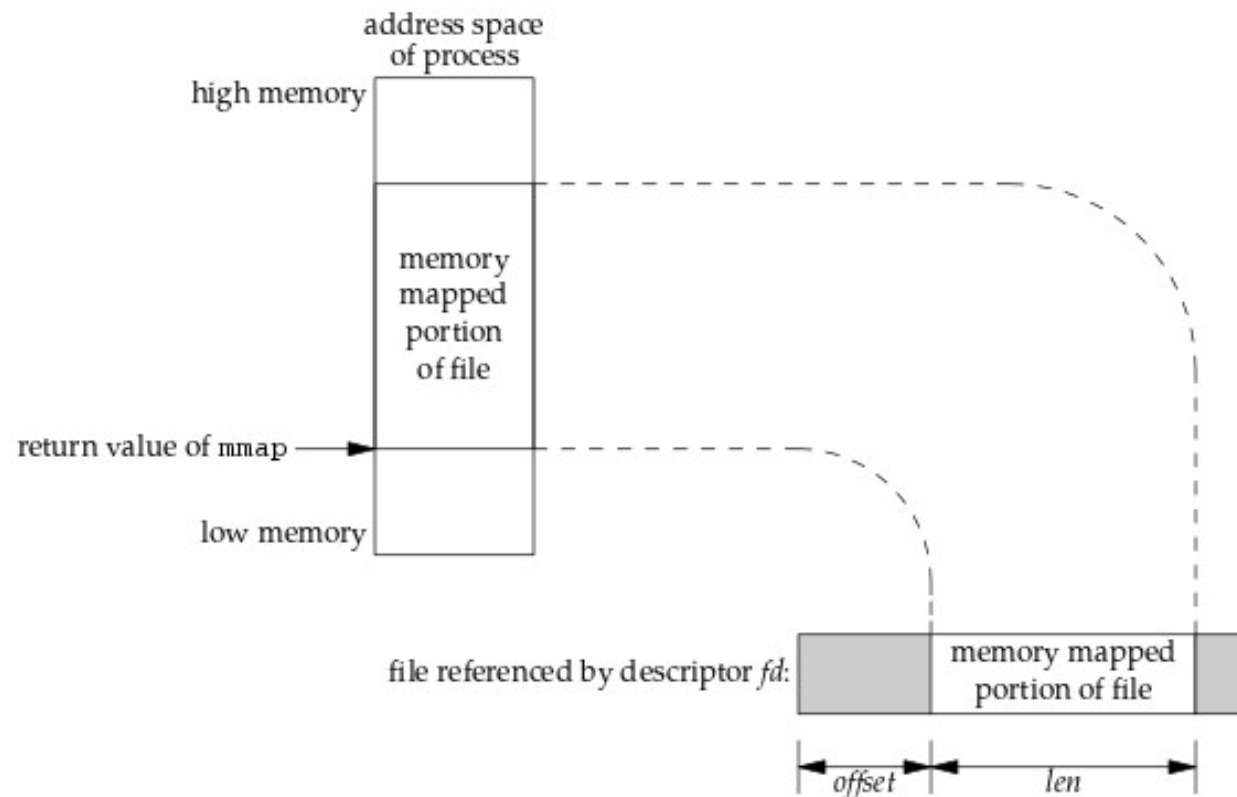
- `int shmget(key_t key, size_t size, int shmflg);`
 - key known by all processes
 - Flags - `IPC_CREAT` | `0666`
- `void *shmat(int shmid, const void *shmaddr, int shmflg);`
 - Shmid – returned by `shmget`
 - shmaddr – `NULL` or other address
 - Shmflg - `SHM_EXEC` `SHM_RDONLY`

System V shared memory

- `char * shm;`
- `key = 5678;`
- `/* Create the segment */`
- `if ((shmid = shmget(key, SHMSZ, IPC_CREAT | 0666)) < 0) {`
- `perror("shmget"); exit(1);`
- `}`
- `/*Now we attach the segment to our data space.*/`
- `if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {`
- `perror("shmat"); exit(1);`
- `}`
 - Repeated on several processes
 - Key must be known
 - Followed by fork
 - Shm value is shared

Memory mapped files

- Access a file context
 - With memory access operations
 - Assignments e accesses



Memory mapped files

- `void *mmap(void * addr, size_t len ,
int prot, int flags,
int fd, off_t offset);`
 - Prot –
 - PROT_READ PROT_WRITE PROT_EXEC PROT_NONE
 - Flags
 - MAP_SHARED MAP_PRIVATE MAP_FIXED
- `int munmap(void * addr, size_t len);`

Memory access

- After mmap a variable contains a pointer to region
 - Programmer can access that memory as a
 - Pointer to variable
 - Vector
- Is it possible to create linked lists in shared memory?
 - No. mmap in different processes returns different addresses

Memory mapped files

- `Int * ptr;`
- `fd = open(argv[1], O_RDWR | O_CREAT, FILE_MODE);`
- `ptr = Mmap(NULL, sizeof(int), PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);`
- `Close(fd);`

Mmap and fork

- Open could be avoided
 - Parent and son share the address to memory
 - Memory is shared among processes
- 4.4BSD introduced anonymous memory sharing
 - Flag – MAP_SHARED|MAP_ANON
 - Fd -1
- ```
ptr = Mmap(NULL, sizeof(int),
 PROT_READ | PROT_WRITE,
 MAP_SHARED | MAP_ANON,
 -1, 0);
```

# POSIX shared memory

- Mmap
  - File system incurs overhead
- Sharing between unrelated processes, without overhead of filesystem I/O
- Intended to be simpler and better than older APIs
- New function to create memory regions

# POSIX shared memory

- Create/opens a shared memory space
  - `fd_mem = shm_open("/myregion", /*region name*/`
  - `O_CREAT | O_RDWR, 0600);`
  - Memory regions are created with size 0
- Assign a size
  - `ftruncate (fd_mem, sizeof(int))`
  - If the object has already been sized by another process, you can get its size with the `fstat` function
- A global region has been created
  - Data stored is kernel persistent
  - But still inaccessible by processes → use `mmap` with `fd_mem`

- `int fd = shm_open(memname,`
- `O_CREAT | O_TRUNC | O_RDWR,`  
`0666);`
- `if (fd == -1)`
- `error_and_die("shm_open");`
- `int r = ftruncate(fd, sizeof(int));`
- `if (r != 0)`
- `error_and_die("ftruncate");`
- `int *v_int = mmap(0, sizeof(int),`
- `PROT_READ | PROT_WRITE, MAP_SHARED,`
- `fd, 0);`



# Shutdown

- Close the shared memory object
  - `close(fd_mem)`
- Unmap the shared memory object:
  - `munmap (pointer, SHM_SIZE);`
  - The address become free to be used.
- At this stage if other processes open and map a memory region
  - They can access data previously written
- To remove permanently the shared memory object:
  - `shm_unlink (SHARED_MEMORY_NAME);`
  - The object is effectively deleted after the last call to `shm_unlink`

# Synchronization on shared memory

- Multiple processes accessing same variable
- Requires synchronization
- 
- Mutexes / condition variables
- Semaphores

# Mutexes across processes

```
shm = (buffer_t *)mmap(NULL, sizeof(buffer_t),
 PROT_READ|PROT_WRITE, MAP_SHARED, zfd, 0);
pthread_mutex_attr_init(&mutex);
pthread_mutexattr_setpshared(&mutex,
 PTHREAD_PROCESS_SHARED);
pthread_mutex_init(&shm->lock, &mutex);
fork()
```

- shm->lock is in shared memory