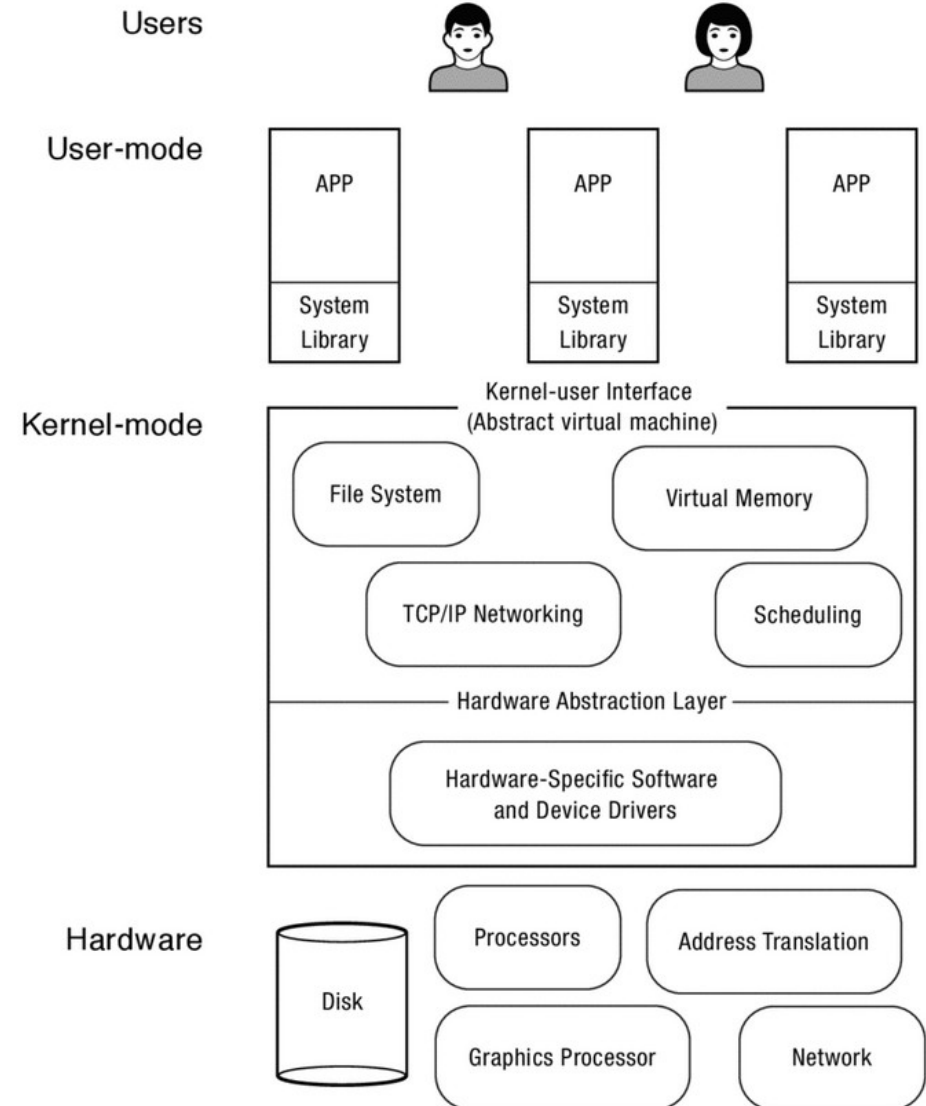
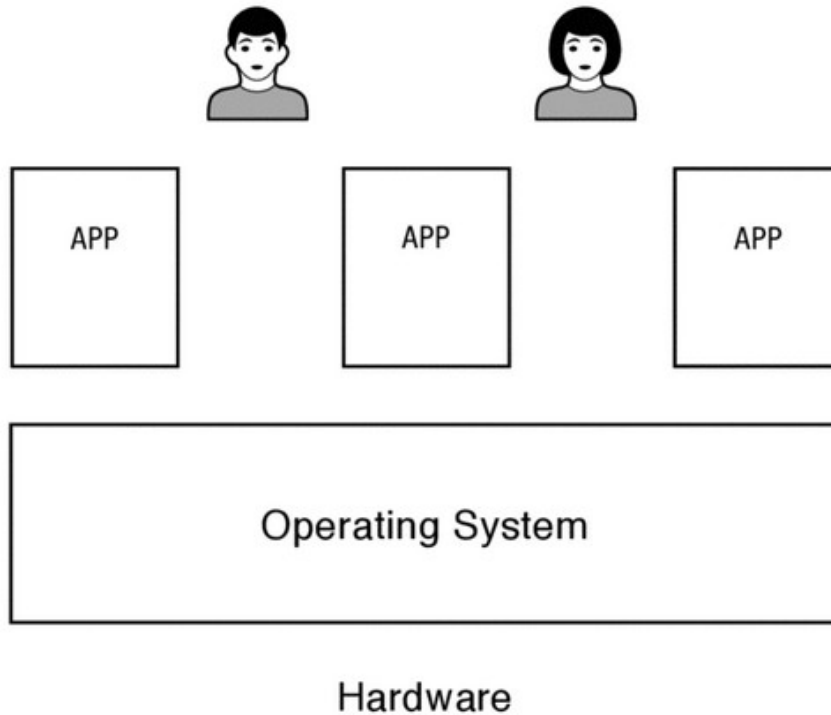


Operating Systems

- Low level software system that
 - manages all applications
 - implements an interface between applications and resources
 - manages available resources
- Resource manager
- Interface
- Virtual Machine

Operating System



Operating System

- OS kernel
 - Code executed in privileged mode
- User space
 - Code executed in non privileged mode
- Service / daemons
 - Application that executed in the background (server)
- Utility programs
 - Application provided by the OS and executed by the user (editor, shell, compiler())
- System calls
 - functions that implement parts of the OS services or utilities
 - can be used inside the kernel
 - Manage and change internal structure.

Operating System Roles

- Referee
 - management of resources shared between application
 - Isolation of applications from each other
 - Decide which applications get which resources
- Illusionist
 - Provides abstraction of physical hardware
 - I/O, available memory, ...
- Glue
 - Provide standard common services
 - I/O (disk, network, ...)
 - Service implementation

Operating System requirements

- Reliability
- Availability
- Security
- Portability
- Performance
- Adoption

Operating System requirements

- Reliability
 - Capacity to work correctly
 - Hide HW faults
 - Reduce SW bugs
 - Handle attacks
- Availability
 - Time the system is working
 - Affected by frequency of failures (MTBF) and time to restore system
- Buggy system that crashes frequently but never loses information
 - Reliable but not available
- Subverted system that appears to be working correctly
 - Available but not reliable

Operating System requirements

- Security
 - Guarantee that operations are not compromised by malicious attacks
 - Privacy
 - guarantee that data is only accessible to authorized users
- SW has bugs that can be exploited
- Administrator can be untrustworthy
- SW developer can be untrustworthy
- OS Design should minimize vulnerabilities
 - Computer virus
 - Downloaded code

Operating System requirements

- Portability
 - Applications are given an abstraction for the HW
 - Resources
 - Access
- Portable abstractions do not change with time
- Abstract virtual machines
 - WIN32
 - POSIX

Operating System requirements




- Performance
 - Efficient use of of available resources
- Overhead
 - added cost of implementing an abstraction
 - Efficiency
 - lack of overhead
- Fairness
 - how to divide resources (memory, CPU) to multiple applications
- Response Time
 - how long a task takes to run
- Throughput
 - rate at which tasks are completed
- Performance predictability

Operating System requirements

- Adoption rate
 - depend on
 - Availability of Applications
 - Availability of hardware
- Open vs closed
- Standardized API
 - more guaranteed of stability
 - more guarantees of portability
- Interoperability

Past, Present Future

- 1945: ENIAC-Electronic Numerical INtegrator And Computer,
 - 19K valves e 1K relays
 - 200 KW
 - 167 m²
- Numeric computations
 - 5000 sums per second or 357 mult per second or 38 divs per second



PDP-7

- First computer to run UNIX
- Introduced in 1965
- Less expensive alternative to the PDP-4
 - price of only U.S.\$72,000
- 18-bit word length
- Standard main memory was 4K words (equivalent to nine kilobytes)
 - 64K words (144 KB)

Motorola 68000


- Developed in the start of 80's
- Used by HP, SUN, DEC, CGI, Apple
- 32-bit
 - instruction set, registers and internal data bus
 - 16-bit external data bus
- Virtual memory
- Frequency < 10Mhz

Speed/Price

| | | CPU | • RAm | Price | MIPS |
|------|--------------------|-------------|----------------|-----------|-----------|
| 1945 | ENIAC | | | 6 936 746 | 0.0000182 |
| 1965 | DEC PDP-8 | | • 200kb | 146 508 | |
| 1983 | IBM PC/XT | i8088 | 16 kB – 256 kB | 12 870 | 0.25 |
| 1984 | Apple Macintosh | 6800 | 128 KB | 6100 | 0.50 |
| 1989 | Sun SS1 | Sparc | 64 MB | 18 598 | 10 |
| 1991 | Dell 433P | i486DX | 32Mb | 5 645 | 11 |
| 1991 | HP 730 | HP-PA7000 | 32 | 28,029 | 78 |
| 2001 | SUN Blade 1000 | UltraSPARC | 1 Gb | 14 477 | 4200 |
| 2001 | | 2xPentium 4 | 1Gb | 1,447 | 2500 |

Past, Present Future

- Processor idle while reading data
- Batch processing
 - Load
 - Run
 - Unload
- While a job was running
 - OS sets I/O devices for next job




Past, Present Future

- Processor idle while reading data
- Multitasking / multiprocessing
 - multiple programs loaded into memory
 - OS selects program to execute
 - Blocking programs (I/O) do not use CPU
- Efficient
 - if queue of tasks is long
 - I/O devices feed processor
- requires program isolation

Past, Present Future

- Processor idle while reading data
- User can interact with the programs
 - multiple display/keyboards
 - make program idle most of the time
- Time-sharing
 - multiple programs loaded into memory
 - Programs from different users
 - OS selects program to execute
 - Blocking programs (I/O) do not use CPU
- Efficient
 - if programs require I/O
- Requires program isolation
- Requires user isolation

- CPU Usage



| |
|------------------|
| Process 1 |
| Process 2 |
| Process 3 |
| Operating System |

- Memory usage

History of UNIX

Past, Present Future

- Modern OS

- Desktop
- Smartphone
- Server
- Embedded
- Virtual machines
- Server clusters
- Cloud

- Future OS

- VLS datacenters
- VLS Multicore systems
- Mobile computing
- Heterogenous systems
- Large Storage
- IoT

The UNIX Time-Sharing System

- Most important achievement
 - demonstrate that powerful operating system for interactive use need not be expensive
 - neither in equipment or in human effort:
- UNIX can run on hardware costing as
 - little as \$40,000,
- the main system software development
 - less than two man years

The UNIX Time-Sharing System

- The PDP-11/45
 - \$40,000
 - 16-bit word (8-bit byte)
 - 144Kbytes of core memory
 - 1M byte fixed-head disk
 - 2.5M bytes on removable disk cartridges
 - console typewriter

The UNIX Time-Sharing System

- Most important job of UNIX is to provide a file system.
 - Ordinary Files
 - Directories
 - Special Files
- Removable File Systems
- Protection
- I/O Calls
 - designed to eliminate the differences between the various devices and styles of access

The UNIX Time-Sharing System

- Processes and Images
 - An image is a computer execution environment. Iti
 - core image (memory)
 - general register values
 - status of open files, current directory, and the like.
- An image is the current state of a pseudo computer.

The UNIX Time-Sharing System

- Processes and Images
 - A process is the execution of an image.
 - While the processor is executing on behalf of a process
 - the image must reside in core
 - Segment organization
 - text segment begins at location 0 in the virtual address space.
 - During execution, this segment is write-protected and a single copy of it is shared among all processes executing the same program.
 - At the first 8K byte boundary above the program text segment
 - begins a non-shared, writable data segment, the size of which may be extended by a systemcall.
 - Starting at the highest address
 - stack segment, which automatically grows down-ward as the hardware's stack pointer fluctuates.

The UNIX Time-Sharing System

- Processes communciation
 - Processes may communicate with related processes through pipes
 - Using he same system read and write calls that are used for file system I/O:
- Traps
 - The PDP-11 hardware detects a number of program faults
 - such as references to nonexistent memory, unimplemented instructions, and odd addresses used where an even address is required.
 - Such faults cause the processor to trap to a system routine
 - Programs which are looping,
 - may be halted by the use of the interrupt signal,

The UNIX Time-Sharing System

- 1969:
 - Bell withdraws from Multics;
 - Thompson starts pondering new os
 - ARPAnet
 - Thompson writes 1st Unix on PDP-7 in one month based on Multics
- 1970
 - Kernighan suggests name 'UNICS' (Uniplexed Information & Computing Service)
 - Thompson writes B language
 - DEC announces PDP-11 [fansite] (will sell 250k)
 - Bell orders \$65k PDP-11/20 w/24kb RAM and 500k disk

The UNIX Time-Sharing System

- 1971
 - Unix rewritten for PDP-11/20
 - os uses 16k, files limited to 64kb
 - **adds pathnames**
 - Unix version 1 written in B
- 1972
 - Ritchie creates C language
 - Unix version 2 written mostly by Thompson in C
- 1973
 - Unix version 3 introduces **pipes text streams** as universal interface
 - multiprogramming added
 - #include and #define added to C
 - Unix kernel rewritten in C
 - Berkeley gets interested in Unix

The UNIX Time-Sharing System

- 1974
 - AT&T decides to supply Unix-source free to academia
- 1976
 - Xerox invents Ethernet
 - Unix licensed by 138 institutions
- 1977
 - Berkeley Software Distribution (Pascal and ex editor)
 - Unix v6 ported for first time
 - Unix clone 'Coherent'

The UNIX Time-Sharing System

- 1978
 - Unix Timesharing System version 7 (UTS)
 - K&R's "The C Programming Language"
 - 'Idris' is Unix clone for PDP-11 purposely incompatible
 - Bill Joy writes **vi** [story]
 - Second Berkeley Software Distribution (2BSD)
 - with **Pascal**, **vi**, **termcap**, **Mail**, **more**, **csh**, **ex**
 - **Berkeley adds virtual memory to Unix 7**

-

The UNIX Time-Sharing System

- 1979
 - **Unix version 7**
 - Bourne shell, awk, lint, make, uucp; find, cpio, expr;
 - large file-systems,
 - unlimited users
 - kernel is still just 40kb
 - Santa Cruz Operation (SCO) founded by Doug and Larry Michels to create Unix ports
 - 1979: Motorola introduces 68000 16-bit CPU
 - IBM prototypes PC using still-scarce 68000
 - Intel introduces 8088 as discount 8-bit version of 8086
 - IBM chooses 8088 for 'short run' PC
 - Unix 32V for VAX [source code]
 - 3BSD is 32-bit port of 2BSD w/ **virtual memory and C shell**

The UNIX Time-Sharing System

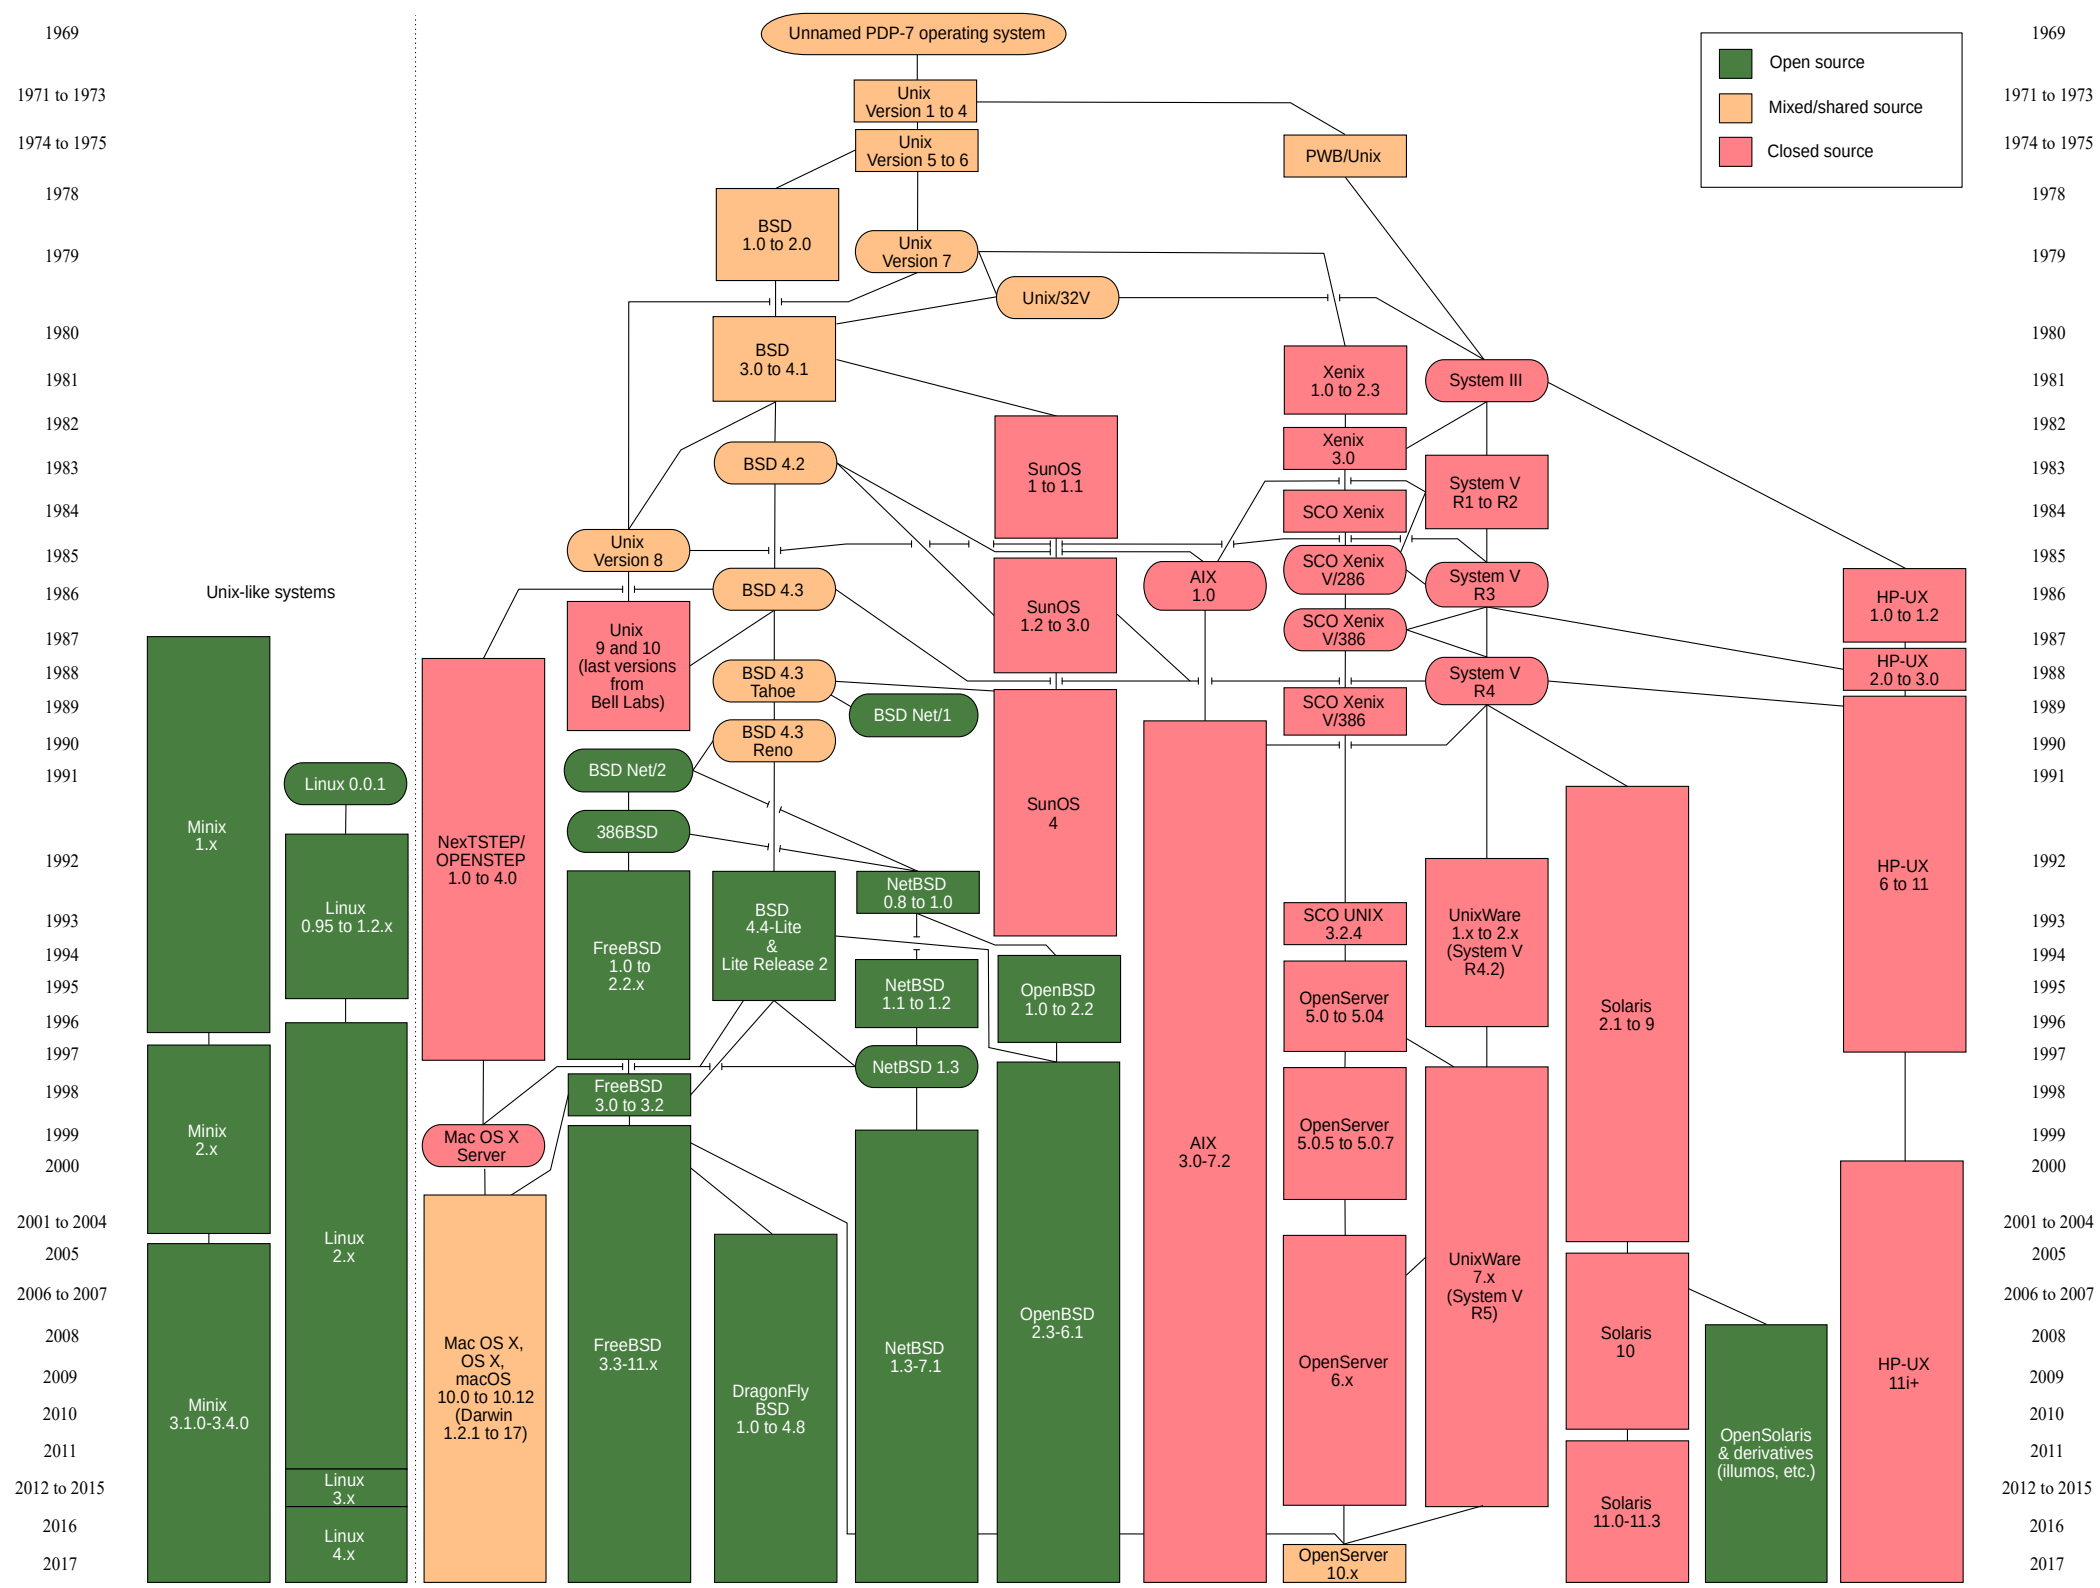
- 1980
 - Bell Labs finally shows interest in BSD Unix
 - Berkeley gets DARPA money because of VAX experience (Comp Systems Research Group)
 - users group for unlicensed (binary-only) users called '**usr**' (later UniForum)
 - Microsoft's 1st-ever os 'Xenix' announced for 8086-et-al,
 - based partly on BSD
 - heavily used at MS for software development and documentation
 - 4BSD incl Franz Lisp, **DARPA enhancements**, curses

The UNIX Time-Sharing System

- 1981
 - '4.1BSD includes speed tweaks and Delivermail
 - no-date: Unix lookalikes include: Onyx, Zeus, Cromix, Omnyx
 - 37yo professor Andy Tanenbaum of Free University in Amsterdam starts work on Minix teaching-system in response to AT&T's stricter license
- 1982
 - Intel intros 80286 with protected memory mode
 - AT&T's 1st commercial Unix, System III
 - AT&T 6300 w/PC Unix
 - no-date: Inix
 - MS Xenix 3.0 will be based on System III but include some 4.1BSD?
 - IBM's version for AT will be called IBM Xenix 1.0;
 - '4.1aBSD' adds TCP/IP and rlogin, etc sendmail?
 - Microsoft promises Xenix ports for PDP-11, 8086, Z8000, and 68000

The UNIX Time-Sharing System

- 1983
 - Apple Lisa runs MS Xenix?
 - DEC abandons PDP-10, MIT shifts to VAX w/Unix
 - MIT launches Project Athena
 - networked workstations to replace Multics
 - 4.2BSD released
 - Stallman announces GNU project
- 1984
 - AT&T divested, Unix becomes commercial product;
 - source code restricted
 - 750 universities have Unix licenses
 - X/Open formed by 5 European computer manufacturers: Bull, ICL, Siemens, Olivetti, and Nixdorf (BISON)
 - MS Xenix System V begins new numbering scheme: IBM calls it IBM Xenix 2.0, SCO calls it SCO Xenix 2.0
 - Earliest SCO copyrights on Xenix 3.x
 - Stallman starts GNU Emacs



The UNIX Time-Sharing System

- 1985
 - Stallman founds Free Software Foundation
 - POSIX standard (Portable Operating System) to reconcile BSD and AT&T variants
 - X-Windows distributed as free
 - X/Open Portability Guide
 - Intel 80386DX can address 4 gigabyte of memory
- 1986
 - 4.3BSD adds speed tweaks
 - Andy Tanenbaum rewrites Unix from scratch
 - for a teaching system on PC/XT/AT/386

The UNIX Time-Sharing System

- 1988
 - AT&T invests in Sun, threatening Unix dominance
 - Open Software Foundation (IBM, DEC, HP, et al)
 - builds on Mach and IBM's AIX; develops Motif
 - NeXT chooses Mach kernel
- 1989
 - open-source components of 4.3BSD released as Networking Release 1
 - AT&T System V v4.0 unifies Xenix, SunOS, 4.3BSD
 - Bill and Lynne Jolitz conceive free 386BSD

The UNIX Time-Sharing System

- 1990
 - 4.3BSD-Reno includes improved virtual memory via Mach
 - Linus takes first C programming class,
 - learns Digital Unix
- 1991
 - Berkeley includes Jolitz's free-but-incomplete 386 BSD on updated Networking Release 2
 - BSD Networking Release 2 includes all but six Unix files
 - linux 0.01 (64kb) 0.02 0.1 1

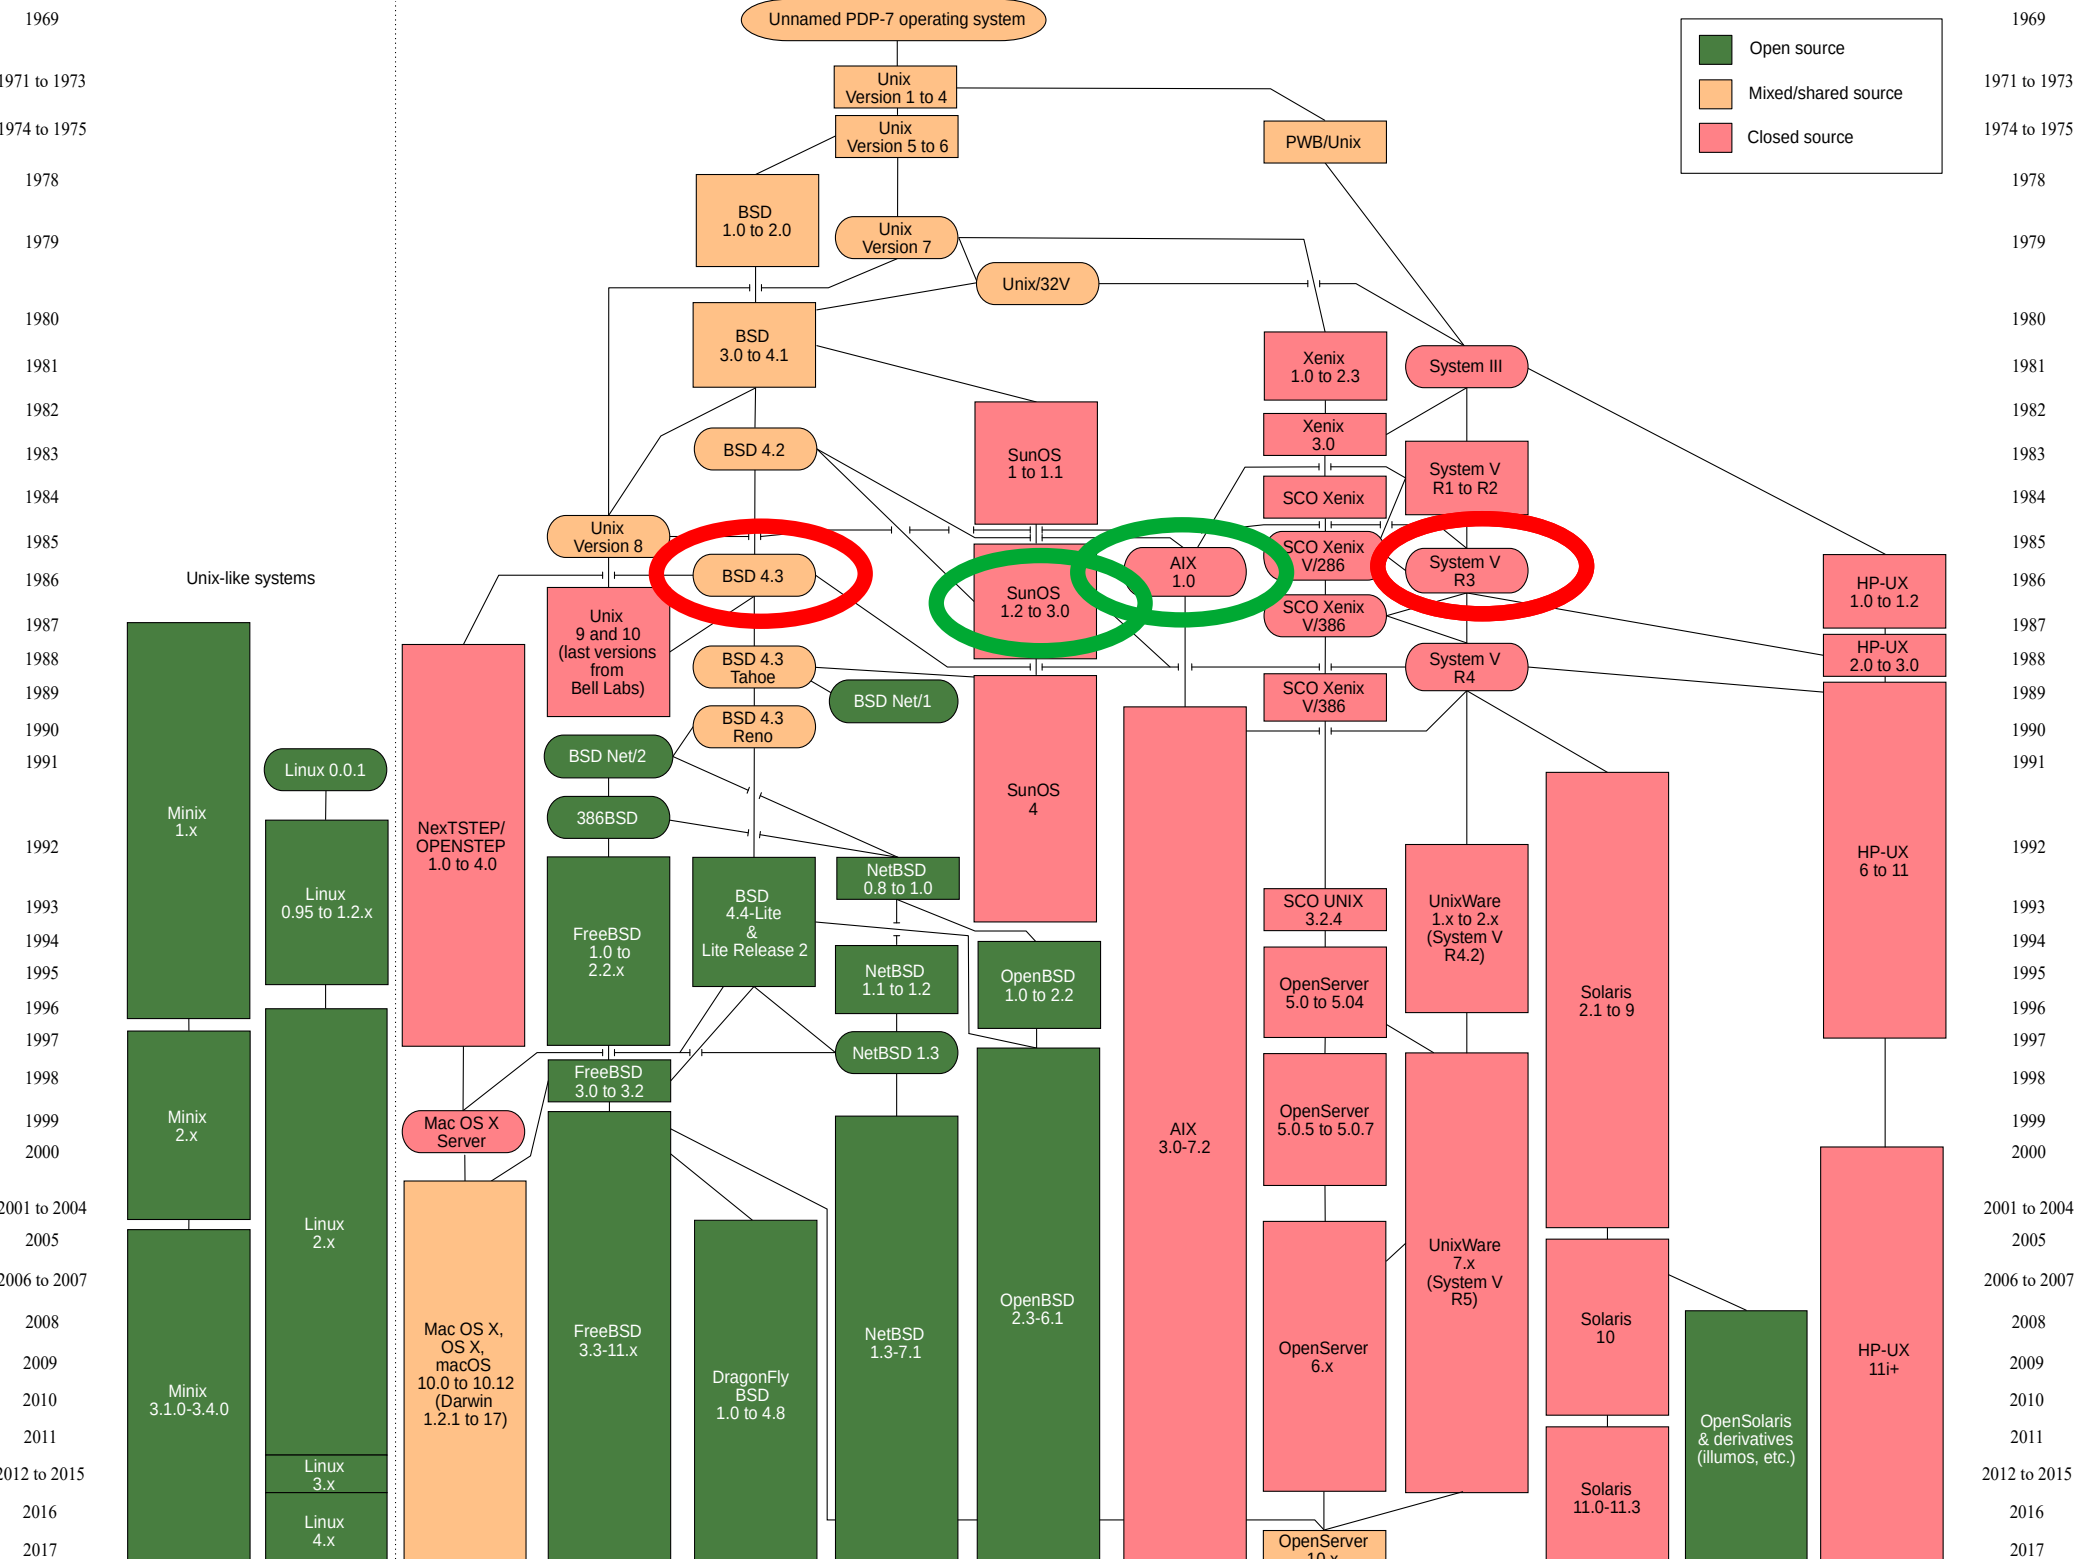
Linux

- In 1991, Linus Torvalds announced the Linux project
 - A free Unix kernel for x86 systems
 - Used Gnu tools from the very beginning
- By 1993, Linux had both internet capability and X capability
 - Just in time for the big internet boom

Standartization

June 1986 USENIX conference in Atlanta

- Many AT&T staff wore buttons which read,
 - “System V: Consider it Standard,”
- and a number of major vendors were promoting products based on System V.
- On the other hand, System V did not yet have TCP/IP networking built innd
 - BSD 4.2 did;
- vendors of engineering workstations were nearly all using BSD, and buttons and posters that said
 - “4.2 > V”



- 1987, AT&T announced that it had purchased a large percentage of Sun Microsystems
 - Sun would receive preferential treatment as AT&T/UNIX Systems Labs developed new software.
- Sun announced that its next system would
 - not be a further extension of SunOS (BSD)
 - but would be derived from AT&T's System V4.
- DEC, in particular, sensed that AT&T was no longer the benign, benevolent progenitor of UNIX

Open Software Foundation

- January 7, 1988 meeting
 - Apollo, DEC, Gould, Hewlett-Packard, Honeywell-Bull, InfoCorp, MIPS, NCR, Silicon Graphics, UniSoft, Unisys
- semi-secret meetings that included
 - HP, IBM, Bull (France), and Nixdorf and Siemens (Germany)
- May 1988,
 - announced the formation of the Open Software Foundation to be dedicated to the production
 - of an operating system, a user interface, a distributed environment,
 - and free cotton candy.

- 1991–92 the worldwide economy worsened.
 - Bull, DEC, IBM, and the computer side of Siemens all lost money.
 - AT&T resold its share of Sun.
- Sun (ATT) had adopted Motif (OSF);
- 1993 USL sold UNIX to Novell,
- OSF abandoned several of its previously announced products
 - shrink-wrapped software and the distributed management environment);
- Bull, Philips, and Siemens withdrew from sponsorship of OSF.

• 1985-93's – FSF and Gnu

- Once you give people something for free, they don't want to have to pay for it
- The Free Software Foundation (1985)
 - Gnu (Gnu's not Unix)
 - Create free versions of popular tools
 - 1986 – gcc, 1987 – most tools
 - Overall goal was to develop a free kernel (It hadn't happened by 1993)
- Squabbling continued and Unix suffered

Free Software Foundation

- Founded by Richard Stallman
 - Wrote original version of gcc and gdb
 - Software should be free, because it should be free
 - As in “free speech, not free beer”
 - Wrote the General Public License (GPL)
 - You are free to do whatever you want as long as the source code goes with it no matter what
 - Controversy: Anything derived from a GPL'd work must itself be GPL'd

Open Source Movement

- Software should be free just because – FSF viewpoint
- Software should be free because free software is better – Open Source Movement viewpoint
 - Every problem can be eliminated if more people look at it

•Linux's Success

- Internet culture and newsgroup postings caused a group of similar minded people to contribute and create a worthwhile kernel
- A competing free attempt had problems
 - The free BSD attempt was mired in a lawsuit (3 files were copied illegally)
 - The Berkeley development group disbanded

OS architectures

- A general purpose OS is composed of:
 - Process manager
 - multiplexes the CPU time between the multiple execution units (processes)
 - memory manager
 - controls, manages and multiplexes the access to physical and virtual memory
 - Inter-process communication
 - Implements and handles mechanism for processes to communicate
 - I/O manager
 - manages communication with peripheral (keyboard/screen, disk, network)
 - User interface
 - command line interpreter
 - GUI

OS organization

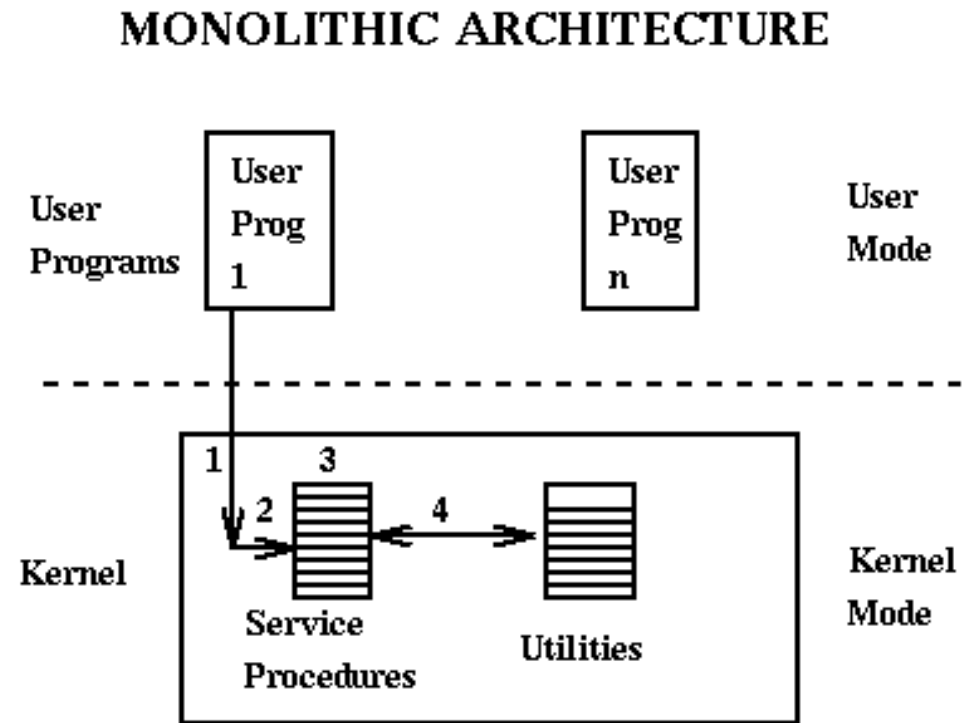
- Layered
- Monolithic
- Micro kernel
- Distributed
- VM based

Monolithic architecture

- OS composed of a single module
 - Although using data abstraction
 - Although using layered approach
- All data and code use same memory space
 - low security mechanisms (one driver can mess other drivers)
 - Difficult to evolve (reboot of system needed)
- Easy to implement
- Low overhead

Monolithic architecture

- DOS
- First Unix versions




1. System call (User->Kernel Mode)
 2. Check parameters
 3. Call service routine
 4. Service Routine call utilities
- Reschedule/Return to user

Layered OS

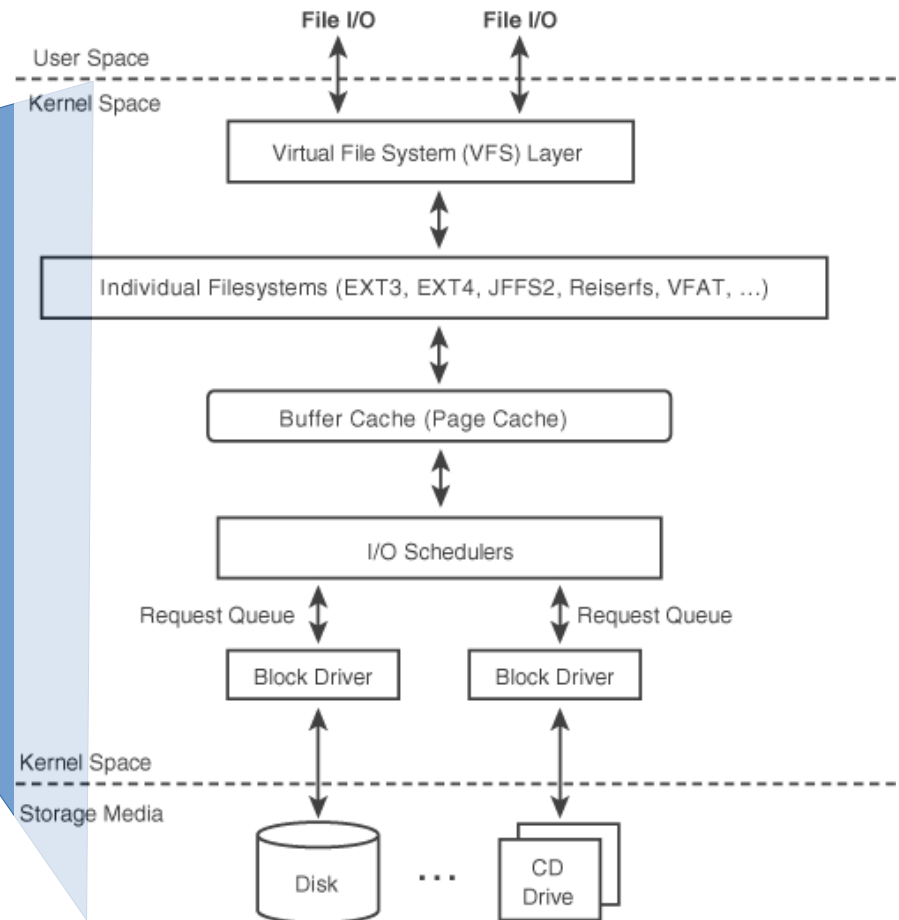
- Components are divided into layers
 - grouping similar components
- Each layer only interacts with:
 - the bottom layer - requesting services
 - to top layer - answering requests
- Higher level layer
 - Applications
- lowest level layer
 - hardware
- Advantaged
 - good structure, well defined interface, ...
- Disadvantages
 - can be slow, may be difficult to define layers.

Hardware
Kernel
Libraries
Shell
Applications



| |
|---------------------------------------|
| OPERATOR |
| USER PROGRAM |
| I/O MANAGEMENT |
| OPERATOR/PROCESS COMMUNICATION |
| MEMORY (MAIN/SECONDARY) MANAGEMENT |
| PROCESSOR ALLOCATION+MULYIPROGRAMMING |
| HARDWARE |

LAYERED SYSTEM (THE System, Dijkstra)



Microkernel

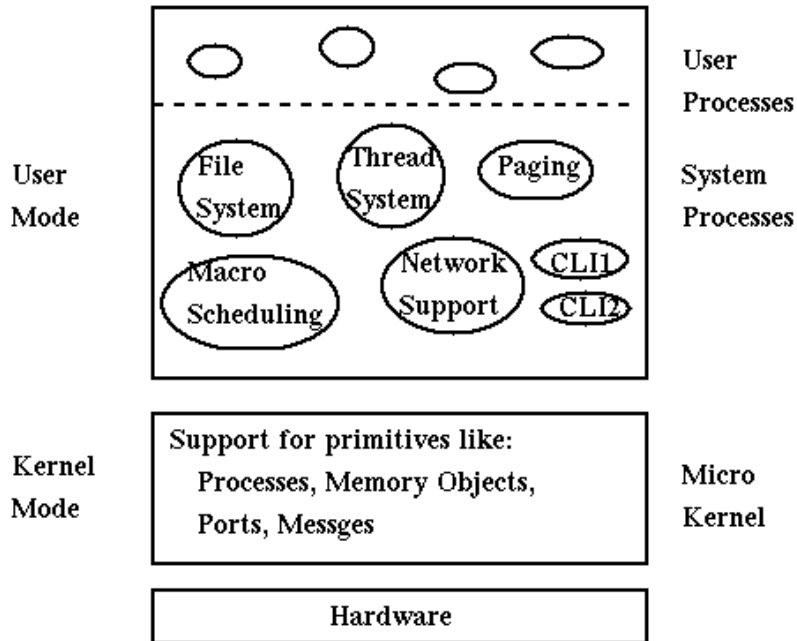
- Removes from kernel as much functionality as possible,
 - limiting the amount of code executed in privileged mode
 - allow easy modifications and extensions
- Most microkernels provide
 - process management
 - memory management
 - message passing between other services
- Security and protection can be enhanced
 - most services are performed in user mode, not kernel mode.
- System expansion can also be easier,
 - only involves adding more system applications, not rebuilding a new kernel.

Microkernel

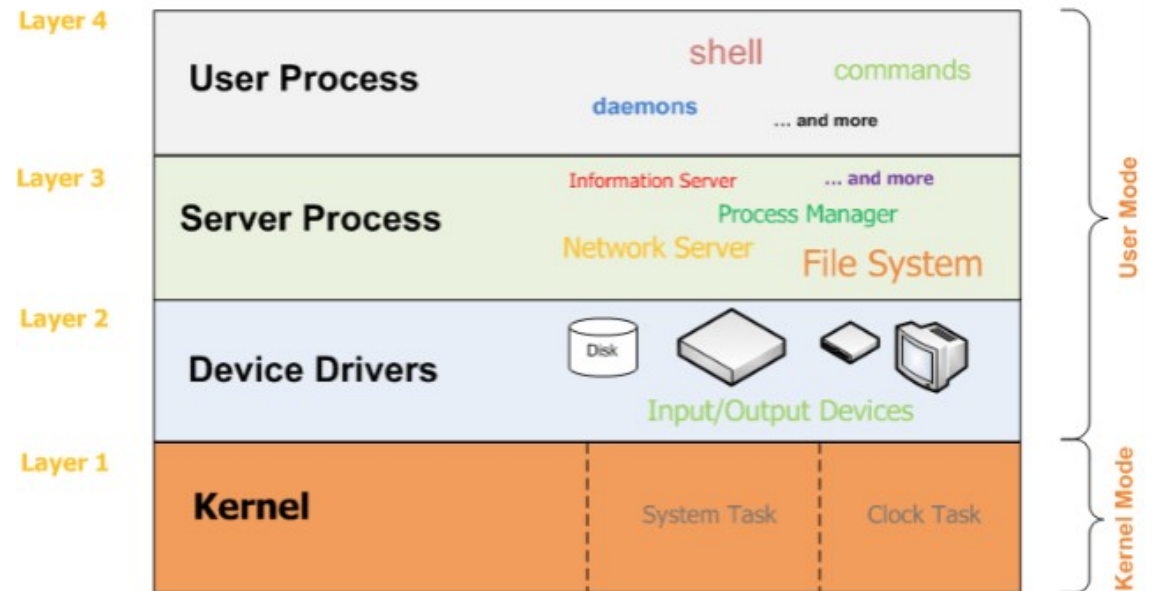
- Windows NT was originally microkernel
 - but suffered from performance problems relative to Windows 95.
 - NT 4.0 improved performance by moving more services into the kernel
- Multiple OSs can be built on top of a micro-kernel
 - Each operating system will make use of different system processes.

Microkernel

Minix Layered Micro Kernel Architecture

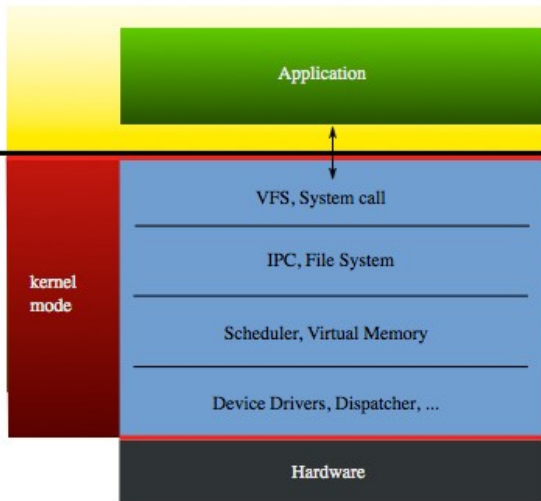


Micro-Kernel Architecture

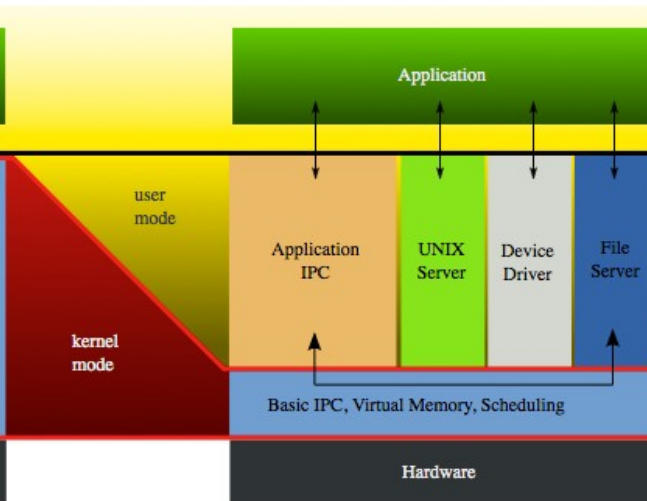


Hibrid kernel

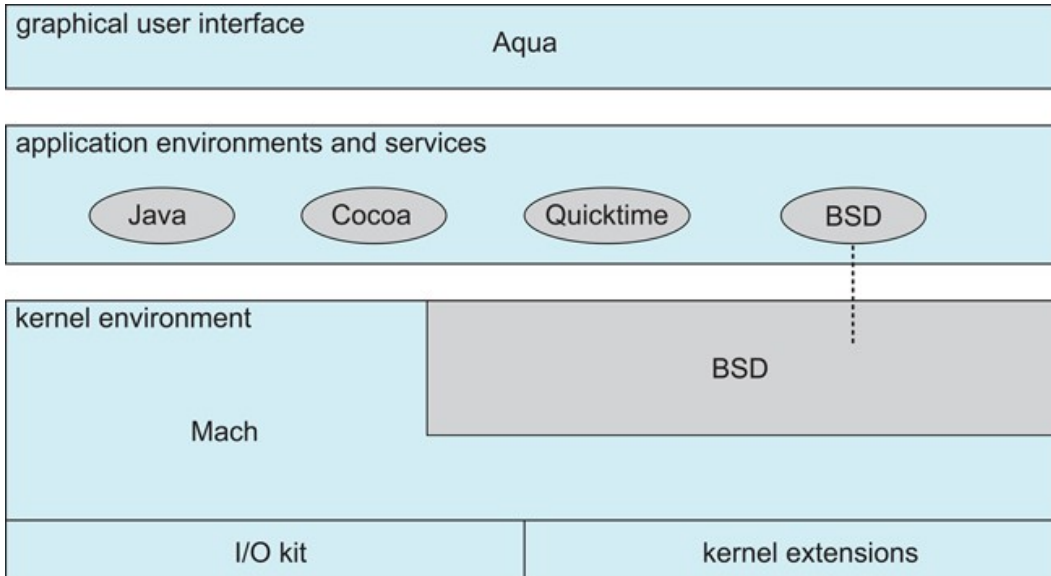
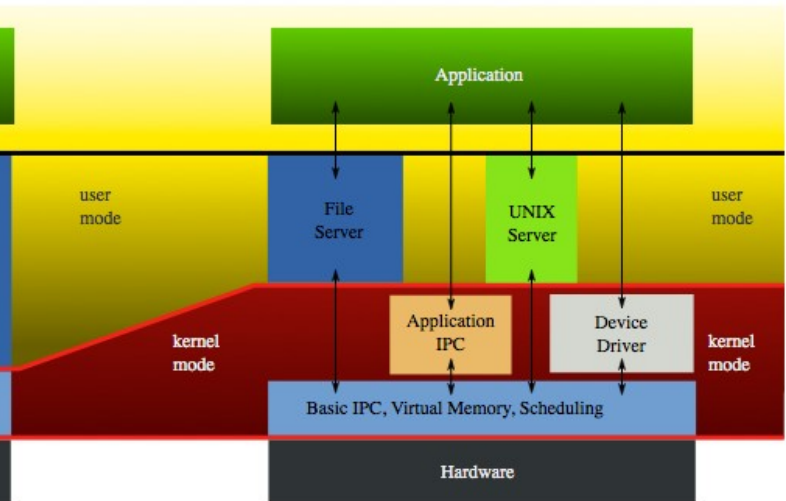
Monolithic Kernel
based Operating System



Microkernel
based Operating System



"Hybrid kernel"
based Operating System



Distributed OS

- Each component/service is a separated process
 - on the same machine
 - on different machines
- Components interactions
 - Messages / Remote procedures
- Distributed File System
- Distributed Memory
- Distributed Processes

Amoeba

- The Bullet Server
 - Used for file storage
- The Directory Server
 - Used for file naming
 - Maps from names to capabilities
- The Replication Server
 - Used for fault tolerance and performance
- The Run Server
 - Run server manages the processor pools
- The Boot Server
 - Ensures that servers are up and running
 - If it discovers that a server has crashed,
 - it attempts to restart it, otherwise selects another processor to provide the service

